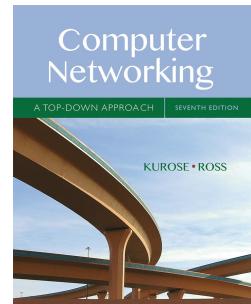


Chapter 2: Application layer

- ❑ 2.1 Principles of network applications
- ❑ 2.2 Web and HTTP
- ❑ 2.3 Electronic Mail
 - ❑ SMTP, POP3, IMAP
- ❑ 2.4 DNS
- ❑ 2.5 P2P applications
- ❑ 2.6 Video streaming and content distribution networks
- ❑ 2.7 Socket programming with UDP and TCP



*Computer Networking:
A Top Down Approach
7th edition
Jim Kurose, Keith Ross
Addison-Wesley
April 2016*

ELEC 331 1

Chapter 2: Application layer

- ❑ 2.1 Principles of network applications
- ❑ 2.2 Web and HTTP
- ❑ 2.3 Electronic Mail
 - ❑ SMTP, POP3, IMAP
- ❑ 2.4 DNS
- ❑ 2.5 P2P applications
- ❑ 2.6 Video streaming and content distribution networks
- ❑ 2.7 Socket programming with UDP and TCP

ELEC 331 2

Chapter 2: Application Layer

Our goals:

- ❑ conceptual, implementation aspects of network application protocols
 - transport layer service models
 - client-server paradigm
 - peer-to-peer paradigm
 - content distribution networks
- ❑ learn about protocols by examining popular application-level protocols
 - HTTP
 - SMTP / POP3 / IMAP
 - DNS
- ❑ programming network applications
 - ❑ socket API

ELEC 331 3

Some network applications

- ❑ e-mail
- ❑ web
- ❑ text messaging
- ❑ remote login
- ❑ P2P file sharing
- ❑ multi-user network games
- ❑ streaming stored video (e.g., YouTube, Netflix)
- ❑ voice over IP (e.g., Skype)
- ❑ real-time video conferencing (e.g., Facetime, Google Hangouts)
- ❑ social networking (e.g., Facebook, Instagram, Twitter, WeChat)
- ❑ location-based mobile apps (e.g., Waze)
- ❑ ...

ELEC 331 4

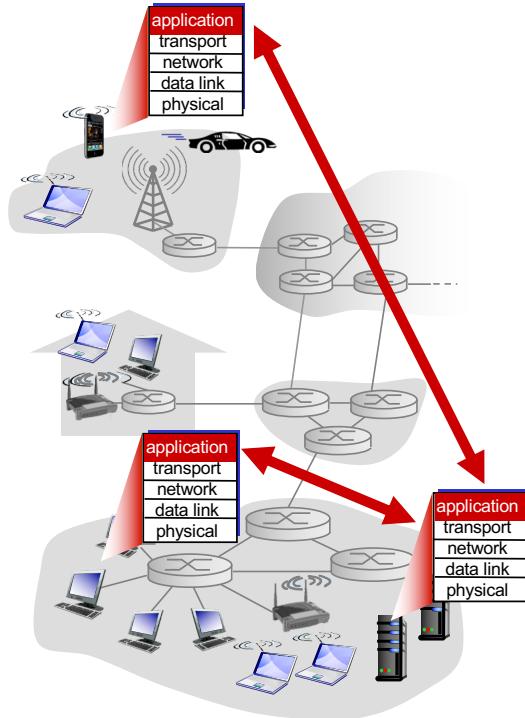
Creating a network app

Write programs that

- run on *different* end systems and communicate over a network
- e.g., Web server software communicates with browser software

No need to write software for network core devices

- Network core devices do not run user's applications
- applications on end systems allow for rapid application development, propagation



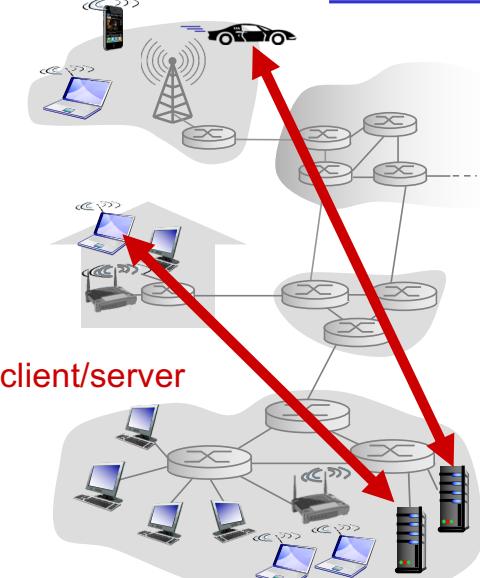
ELEC 331 5

Application architectures

- Dictate how the application is organized over various end systems
- Designed by the application developer
- Possible structure of applications:
 1. Client-server architecture
 - o Including data centers / cloud computing
 2. Peer-to-peer (P2P) architecture

ELEC 331 6

Client-server architecture



Server Farm: cluster of servers

Search engine (Google, Bing), e-commerce (Amazon, eBay), webmail (Gmail), social networking (Facebook, Twitter), video (YouTube)

Server:

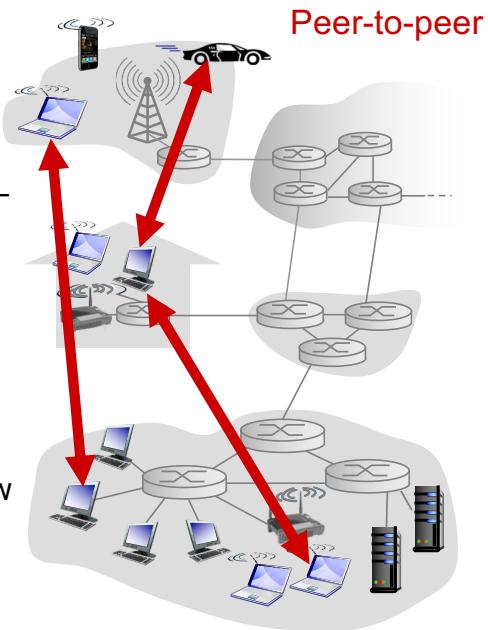
- ❑ always-on host
- ❑ permanent IP address
- ❑ server farms (data center) for scaling

Clients:

- ❑ communicate with server
- ❑ may be intermittently connected
- ❑ may have dynamic IP addresses
- ❑ do not communicate directly with each other

P2P architecture

- ❑ no always on server
- ❑ arbitrary end systems directly communicate
- ❑ e.g., file sharing (e.g., BitTorrent), peer-assisted download acceleration (e.g., Xunlei), Internet telephony and video conference (e.g., Skype)
- ❑ peers request service from other peers, provide service in return to other peers
 - **self scalability** – new peers bring new service capacity, as well as new service demands
- ❑ peers are intermittently connected and may change IP addresses
 - complex management



Processes communicating

Process: program running within a host

- ❑ within same host, two processes communicate using **inter-process communication** (defined by OS)
- ❑ processes in different hosts (with potentially different OS) communicate by exchanging **messages**

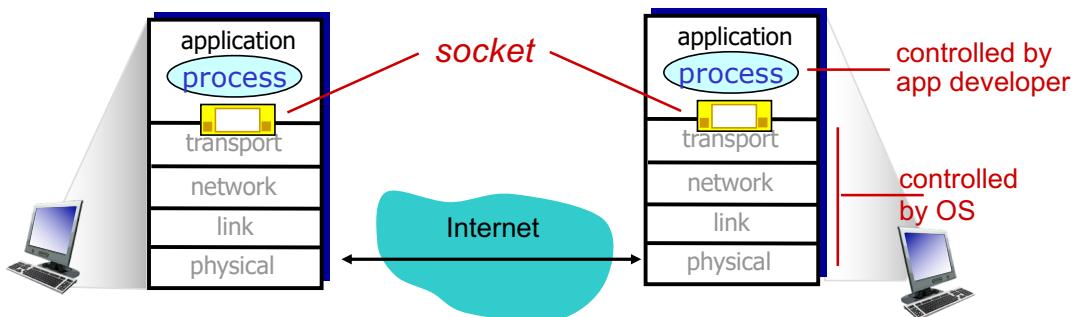
Client process: process that initiates communication

Server process: process that waits to be contacted

- ☞ Note: applications with P2P architectures have client processes & server processes

Sockets

- ❑ process sends/receives messages to/from its **socket**
- ❑ **Process** is analogous to a **house** and **socket** is analogous to **door**
 - sending process shoves message out its door
 - sending process relies on transport infrastructure on other side of door which brings message to socket at receiving process



Addressing processes

- ❑ For a process to receive messages, it must have an identifier
- ❑ A host has a unique 32-bit IP address
- ❑ Ex: use ipconfig (Windows) or ifconfig (Apple) from command prompt to get your IP address
- ❑ Q: does the IP address of the host on which the process runs suffice for identifying the process?
- ❑ Answer: No, many processes can be running on same host
- ❑ Identifier includes both the IP address and **port numbers** associated with the process on the host.
- ❑ Example port numbers:
 - ❑ HTTP server: 80
 - ❑ Mail server using SMTP: 25
- ❑ to send HTTP message to gaia.cs.umass.edu web server:
 - ❑ IP address: 128.119.245.12
 - ❑ Port number: 80
- ❑ More on this later

ELEC 331 11

What transport service does an app need?

Reliable Data Transfer

- ❑ Some apps (e.g., file transfer, web, financial apps) require data be delivered correctly and completely to the destination
- ❑ **Loss-tolerant apps:** Some apps (e.g., real-time audio /video) can tolerate some loss

Timing

- ❑ Some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

Throughput

- r **Bandwidth sensitive apps:** some apps (e.g., multimedia) require minimum amount of bandwidth to be **effective**
- r Other **elastic** apps (e.g., e-mail, web, file transfer) make use of whatever bandwidth they can get

Security

- r Confidentiality (encrypt data), data integrity, end-point authentication

ELEC 331 12

Transport service requirements of common apps

Application	Data loss	Bandwidth	Time Sensitive
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic (few kbps)	no
Internet telephony	loss-tolerant	audio: 5kbps-1Mbps	yes, 100's
video conferencing		video:10kbps-5Mbps	msec
Streaming stored audio/video	loss-tolerant	same as above	yes, few secs
interactive games	loss-tolerant	few kbps–10 kbps	yes, 100's msec
instant messaging	no loss	elastic	yes and no

ELEC 331 13

Internet transport protocols services

TCP service:

- connection-oriented*: setup required between client and server processes (**handshaking**)
- reliable data transfer* between sending and receiving process
- flow control*: sender won't overwhelm receiver
- congestion control*: throttle sender when network overloaded
- does not provide*: timing, min throughput guarantees

UDP service:

- connectionless*: no handshaking
 - unreliable data transfer* between sending and receiving process
 - does not provide: flow control and congestion control
 - can pump data into network layer at any rate it pleases
- Q:** why bother? Why is there a UDP?

ELEC 331 14

Internet apps: Application, transport protocols

Application	Application layer protocol	Underlying Transport Protocol
e-mail	SMTP [RFC 5321]	TCP
remote terminal access	Telnet [RFC 854]	TCP
Web	HTTP [RFC 2616]	TCP
file transfer	FTP [RFC 959]	TCP
streaming multimedia	HTTP (e.g., YouTube)	TCP
Internet telephony	SIP [RFC 3261], RTP [RFC 3550], or proprietary (e.g., Skype)	UDP or TCP

ELEC 331 15

Securing TCP

TCP and UDP

- no encryption
- cleartext (i.e., unencrypted) passwords sent into socket traverse Internet in cleartext

SSL socket API

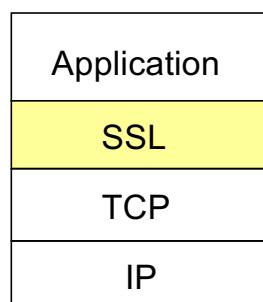
- cleartext passwords sent into socket traverse Internet encrypted
- see Chapter 8

SSL (Secure Sockets Layer)

- provides encrypted TCP connection
- data integrity
- end-point authentication

SSL at application layer

- Apps use SSL libraries, which "talk" to TCP



application with SSL

ELEC 331 16

Application layer protocol defines

- ❑ **Types** of messages exchanged, e.g., request & response messages
 - ❑ **Syntax** of message types: what fields in messages & how fields are delineated
 - ❑ **Semantics** of the fields, i.e., meaning of information in fields
 - ❑ **Rules** for when and how processes send & respond to messages
- Open protocols:**
- ❑ defined in IETF Request for Comments (RFCs)
 - ❑ allow for interoperability
 - ❑ e.g., HTTP, SMTP
- Proprietary protocols:**
- ❑ e.g., Skype

ELEC 331 17

Chapter 2: Application layer

- ❑ 2.1 Principles of network applications
 - ❑ app architectures
 - ❑ app requirements
- ❑ 2.2 Web and HTTP
- ❑ 2.3 Electronic Mail
 - ❑ SMTP, POP3, IMAP
- ❑ 2.4 DNS
- ❑ 2.5 P2P applications
- ❑ 2.6 Video streaming and content distribution networks
- ❑ 2.7 Socket programming with UDP and TCP

ELEC 331 18

Web and HTTP

First, a review

- Web page consists of objects
- Object can be HTML file, JPEG image, Java applet, video clip...
- Web page consists of base HTML file which includes several referenced objects
- Each object is addressable by a URL
- Example URL:

www.someSchool.edu/someDept/pic.gif



ELEC 331 19

HTTP (HyperText Transfer Protocol) overview

- Web's application layer protocol
- client/server model
 - *client*: browser that requests, receives, (using HTTP protocol) "displays" Web objects
 - *server*: Web server sends (using HTTP protocol) objects in response to requests



ELEC 331 20

HTTP overview (continued)

Uses TCP:

- ❑ client initiates TCP connection (creates socket) to server, port 80
- ❑ server accepts TCP connection from client
- ❑ HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- ❑ TCP connection closed

HTTP is “stateless”

- ❑ server maintains no information about past client requests

aside

Protocols that maintain “state” are complex!

- past history (state) must be maintained
- if server/client crashes, their views of “state” may be inconsistent, must be reconciled

HTTP connections

Non-persistent HTTP

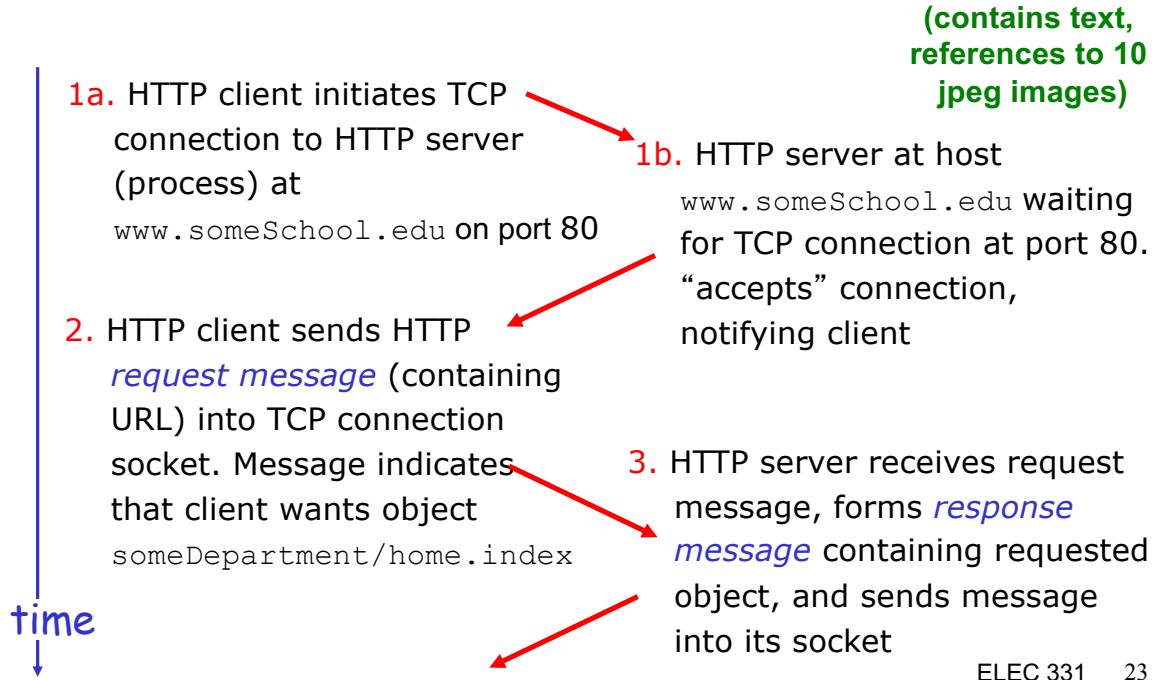
- ❑ At most one object is sent over a TCP connection.
- ❑ Connection is then closed.
- ❑ Downloading multiple objects required multiple connections.

Persistent HTTP

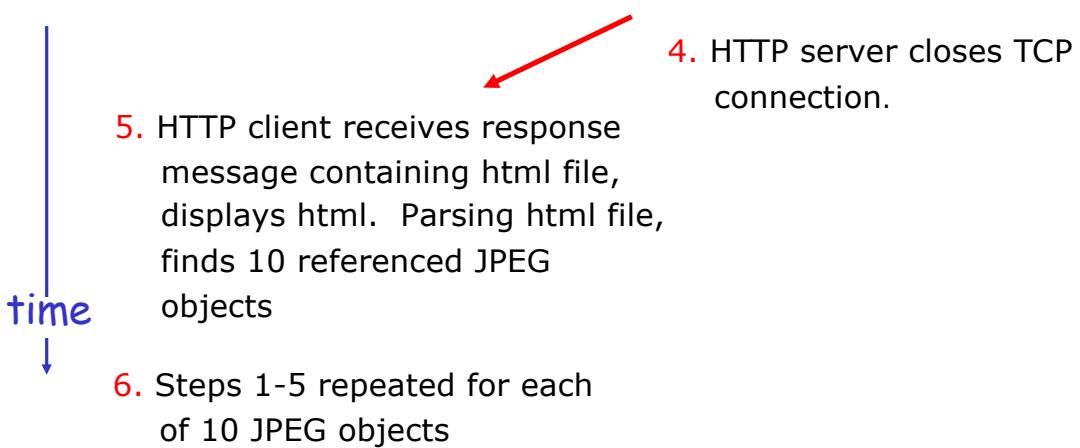
- ❑ Multiple objects can be sent over single TCP connection between client and server.

HTTP with non-persistent connection

Suppose user enters URL `www.someSchool.edu/someDepartment/home.index`



Non-persistent HTTP (cont.)



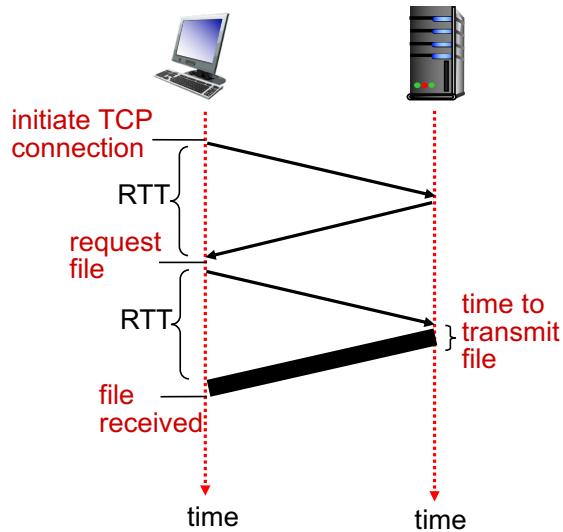
Response time modeling

Definition of RTT: time to send a small packet to travel from client to server and back.

Response time:

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- file transmission time

$$\text{total} = \text{2RTT} + \text{transmission time}$$



ELEC 331 25

Persistent HTTP

Non-persistent HTTP issues

- requires 2 RTTs per object
- OS must work and allocate host resources for each TCP connection
- but browsers often open parallel TCP connections to fetch referenced objects

Persistent HTTP

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server are sent over connection
- client sends requests as soon as it encounters a referenced object (**pipelining**)
- as little as one RTT for all the referenced objects

ELEC 331 26

HTTP request message

- two types of HTTP messages: *request, response*
- **HTTP request message:**
 - ASCII (human-readable format)

request line
(GET, POST,
HEAD, PUT,
DELETE commands)

header lines

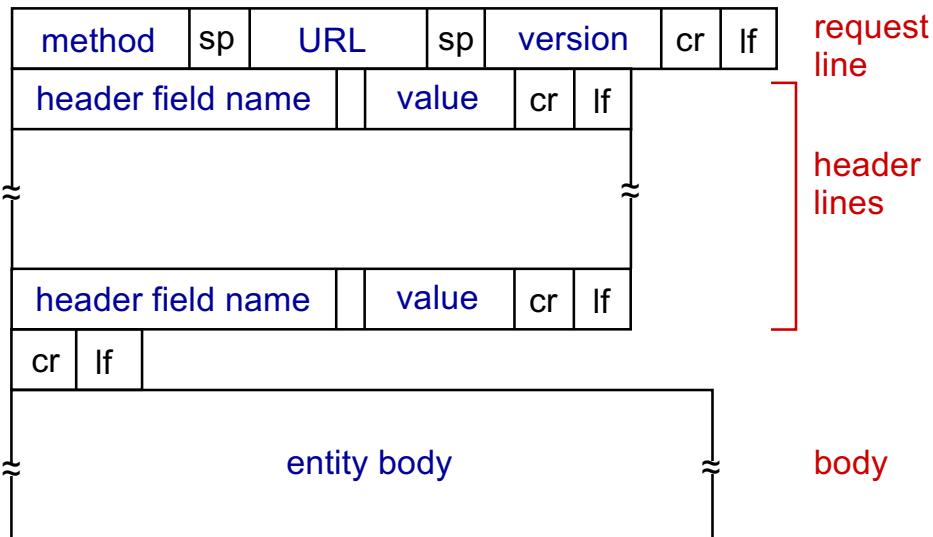
carriage return,
line feed at start
of line indicates
end of header lines

```
GET /index.html HTTP/1.1\r\n
Host: www-net.cs.umass.edu\r\n
User-Agent: Firefox/3.6.10\r\n
Accept: text/html,application/xhtml+xml\r\n
Accept-Language: en-us,en;q=0.5\r\n
Accept-Encoding: gzip,deflate\r\n
Accept-Charset: ISO-8859-1,utf-8;q=0.7\r\n
Keep-Alive: 115\r\n
Connection: keep-alive\r\n\r\n
```

* Check out the online interactive exercises for more examples: http://gaia.cs.umass.edu/kurose_ross/interactive/

ELEC 331 27

HTTP request message: general format



ELEC 331 28

Uploading form input

Post method:

- ❑ Web page often includes form input
- ❑ Input is uploaded to server in entity body

URL method:

- ❑ Uses GET method
- ❑ Input is uploaded in URL field of request line:

www.somesite.com/animalsearch?monkeys&banana

ELEC 331 29

Method types

HTTP/1.0

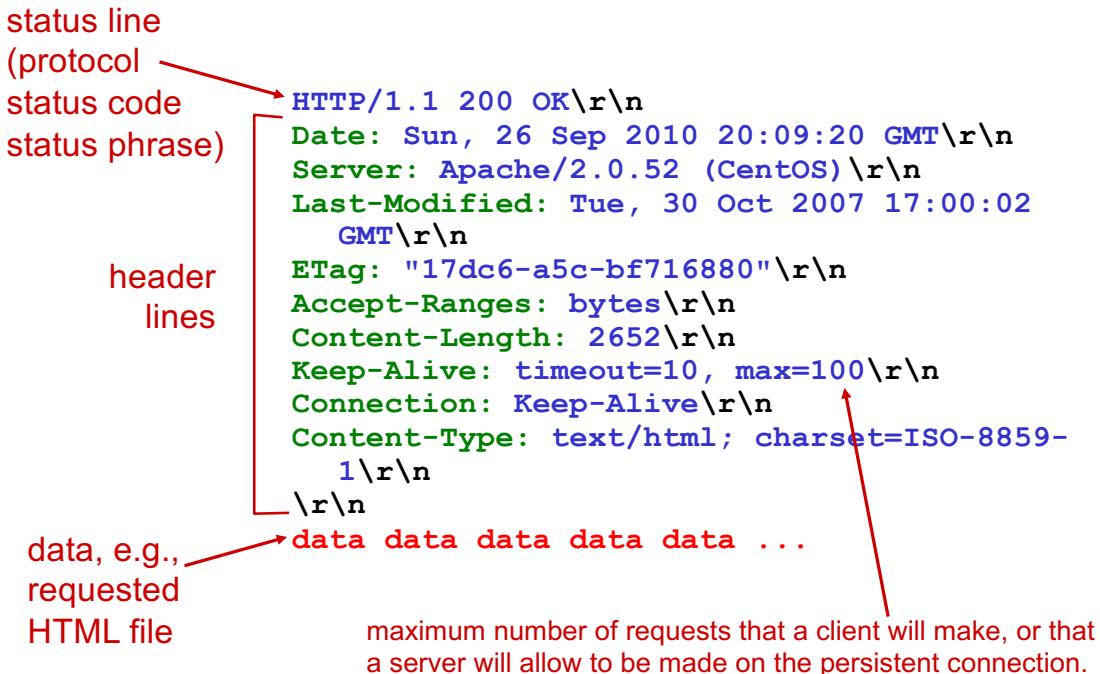
- ❑ GET
- ❑ POST
- ❑ HEAD
 - asks server to leave requested object out of response
 - for debugging

HTTP/1.1

- ❑ GET, POST, HEAD
- ❑ PUT
 - uploads file in entity body to path specified in URL field
 - for web publishing
- ❑ DELETE
 - deletes file specified in the URL field
- ❑ ... and a few others

ELEC 331 30

HTTP response message



* Check out the online interactive exercises for more examples: http://gaia.cs.umass.edu/kurose_ross/interactive/

ELEC 331 31

HTTP response status codes

Status code appears in 1st line in server-to-client response message.

A few sample codes:

200 OK

- ❑ request succeeded, requested object later in this message

301 Moved Permanently or Permanent Redirect

- ❑ requested object moved, new location specified later in this message (`Location: header`)
- ❑ Can be used to support [URL shortening](#). E.g., "http://en.wikipedia.org/wiki/URL_shortening" can be shortened to <http://bit.ly/urlwiki>

400 Bad Request

- ❑ request message not understood by server

404 Not Found

- ❑ requested document not found on this server

505 HTTP Version Not Supported

ELEC 331 32

Trying out HTTP (client side) for yourself

1. Telnet to your favorite Web server:

```
telnet gaia.cs.umass.edu 80
```

Opens TCP connection to port 80
(default HTTP server port) at gaia.cs.umass.edu
Anything typed in sent
to port 80 at gaia.cs.umass.edu

2. Type in a GET HTTP request:

```
GET /kurose_ross/interactive/index.php HTTP/1.1  
Host: gaia.cs.umass.edu
```

By typing this in (hit carriage return twice), you send this minimal (but complete) GET request to HTTP server

3. Look at response message sent by HTTP server!

(or use Wireshark to look at captured HTTP request/response)

4. Try again. This time replace GET with HEAD.

User-server state: cookies

Many major Web sites use
cookies

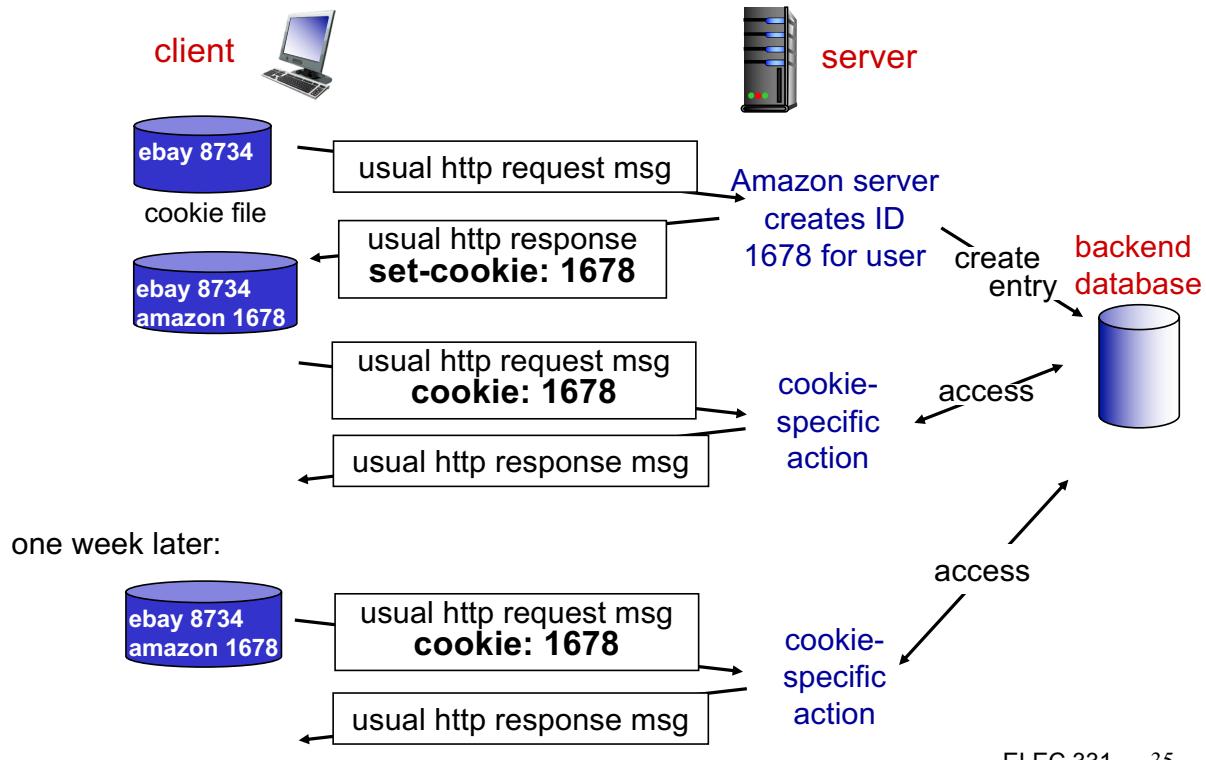
Example:

Four components:

- 1) cookie header line in the HTTP response message
- 2) cookie header line in HTTP request message
- 3) cookie file kept on user's host and managed by user's browser
- 4) backend database at Web site

- Susan accesses Internet always from same PC
- She visits a specific e-commerce site for first time
- When initial HTTP request arrives at site, site creates a unique ID and creates an entry in backend database for ID

Cookies: keeping “state” (cont.)



ELEC 331 35

Cookies (continued)

What cookies can be used for:

- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)

How to keep “state”:

- protocol endpoints: maintain state at sender/receiver over multiple transactions
- cookies: http messages carry state

aside
Cookies and privacy:

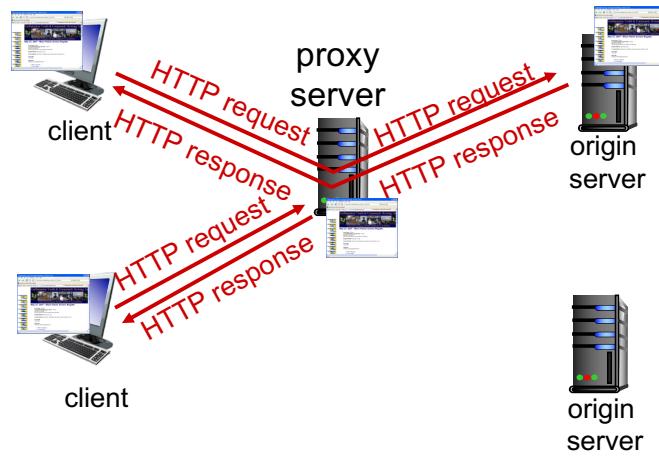
- cookies permit sites to learn a lot about you
- you may supply name and e-mail to sites
- search engines use redirection & cookies to learn yet more
- advertising companies obtain info across sites

ELEC 331 36

Web caches (proxy server)

Goal: satisfy client request without involving origin server

- ❑ user sets browser: Web accesses via cache
- ❑ browser sends all HTTP requests to cache
 - object in cache: cache returns object
 - else cache requests object from origin server, then returns object to client



More about Web caching

- ❑ Cache acts as both client and server
 - server for original requesting client
 - client to origin server
- ❑ Typically cache is installed by ISP (university, company, residential ISP)

Why Web caching?

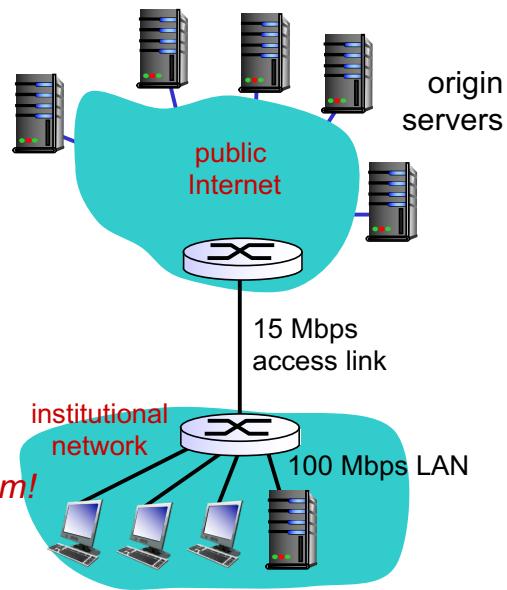
- ❑ Reduce response time for client request.
- ❑ Reduce traffic on an institution's access link.
- ❑ Internet dense with caches enables “poor” content providers to effectively deliver content (but so does P2P file sharing)

Assumptions

- ❑ avg object size: 1 Mbits
- ❑ avg request rate from browsers to origin servers: 15 requests per sec
- ❑ avg arriving data rate on access link: 15 Mbps
- ❑ HTTP request msg negligibly small
- ❑ access link rate: 15 Mbps
- ❑ RTT between router on Internet side and any origin server: 2 sec

Consequences

- ❑ traffic intensity on LAN: 0.15
- ❑ traffic intensity on access link = **1** *problem!*
- ❑ total delay = Internet delay + access delay + LAN delay
= 2 sec + minutes + msecs

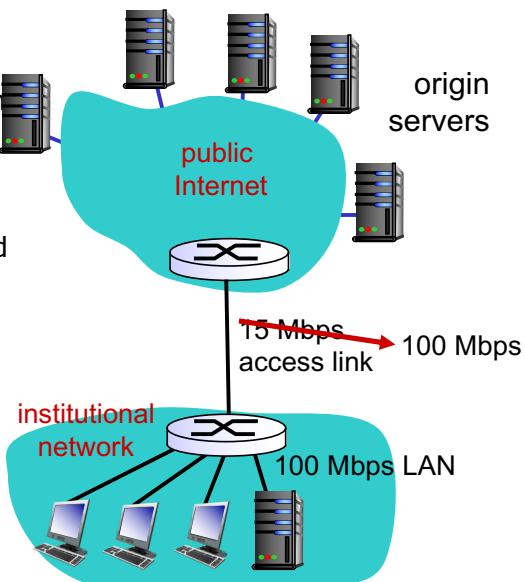


ELEC 331 39

Caching example fatter access link

Possible solution

- ❑ avg object size: 1 Mbits
- ❑ avg request rate from browsers to origin servers: 15 requests/sec
- ❑ avg arrival rate on access link: 15 Mbps
- ❑ RTT between router on Internet side and any origin server: 2 sec
- ❑ access link rate: ~~15 Mbps~~ → **100 Mbps**



Consequences

- ❑ traffic intensity on LAN: 0.15
- ❑ traffic intensity on access link = ~~1~~ → **0.15**
- ❑ total delay = Internet delay + access delay + LAN delay
= 2 sec + ~~minutes~~ + msecs

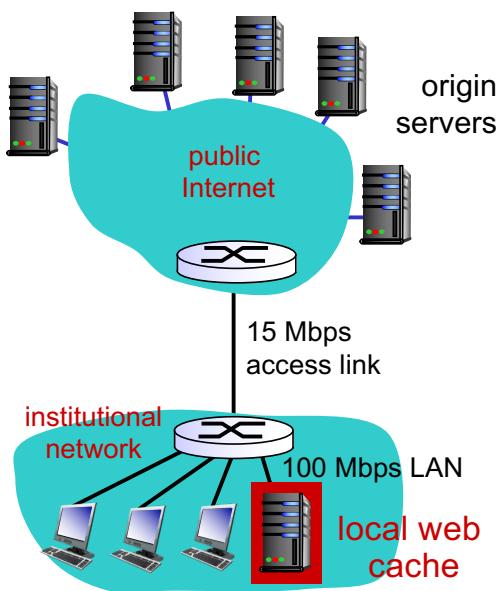
Cost: increased access link speed (not cheap!)

ELEC 331 40

Caching example install local cache

Calculating access link utilization, delay with cache:

- ❑ suppose **hit rate** is 0.4
 - 40% requests satisfied at cache, 60% requests satisfied at origin
- ❑ access link utilization:
 - 60% of requests use access link
- ❑ data rate to browsers over access link = $0.6 \times 15 \text{ Mbps} = 9 \text{ Mbps}$
 - Traffic intensity on access link = 0.6
- ❑ total delay
 - $= 0.6 \times (\text{delay from origin servers}) + 0.4 \times (\text{delay when satisfied at cache})$
 - $= 0.6 (2.01) + 0.4 (\sim \text{msecs})$
 - $\approx 1.2 \text{ secs}$
 - less than with 150 Mbps link (and cheaper too!)



ELEC 331 41

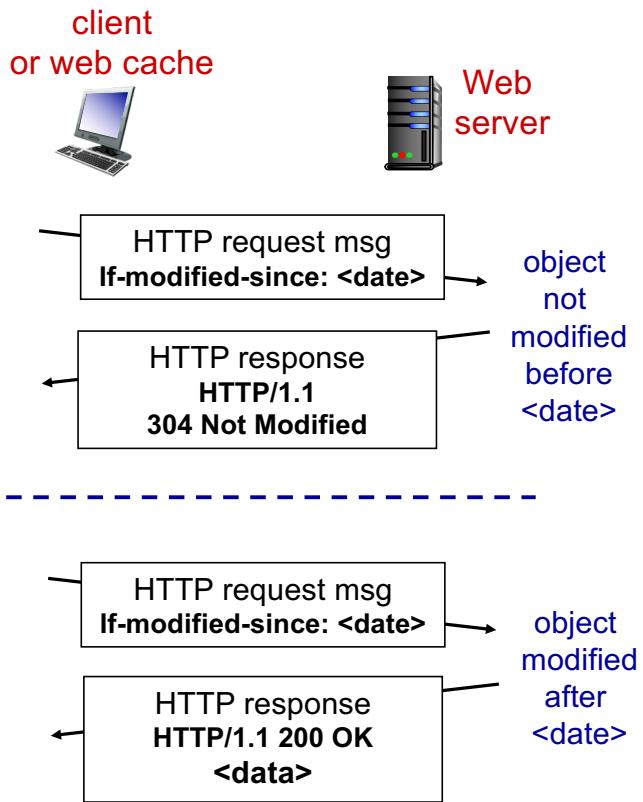
Conditional GET

- ❑ **Goal:** don't send object if cache has up-to-date cached version
 - no object transmission delay
 - lower link utilization
- ❑ **cache:** specify date of cached copy in HTTP request

If-modified-since:
 `<date>`

- ❑ **server:** response contains no object if cached copy is up-to-date:

`HTTP/1.1 304 Not Modified`



ELEC 331 42

Final words on caching

Cache directives: Must be obeyed by all caching mechanisms along the request/response chain. Specify whether or not an object can be cached. How long it can be cached.

Server

- ❑ **Cache-Control** header: Tells all caching mechanisms from server to client whether they may cache this object.. E.g., Cache-Control: : max-age=3600
- ❑ **Expires** header: date/time at which contents of object are considered stale. Expires: Tue, 23 Sep 2013 16:00:00 GMT
- ❑ **Etag (Entity Tag)** header: used to determine if a cached version of the requested object is identical to the current version of the object on the server. If object's content has changed, a new and different Etag is assigned. ETag: "1269c72-44d6-48f50ca515ed2"

Client

- ❑ **Cache-Control** header: Used to specify directives that MUST be obeyed by all caching mechanisms along the request/response chain. E.g., Cache-Control: no-cache
- ❑ **If-none-match** header. Allows a 304 Not Modified to be returned if content is unchanged. E.g., If-None-Match: "284d8af7ad3"

Chapter 2: Application layer

- ❑ 2.1 Principles of network applications
- ❑ 2.2 Web and HTTP
- ❑ 2.3 Electronic Mail
 - ❑ SMTP, POP3, IMAP
- ❑ 2.4 DNS
- ❑ 2.5 P2P applications
- ❑ 2.6 Video streaming and content distribution networks
- ❑ 2.7 Socket programming with UDP and TCP

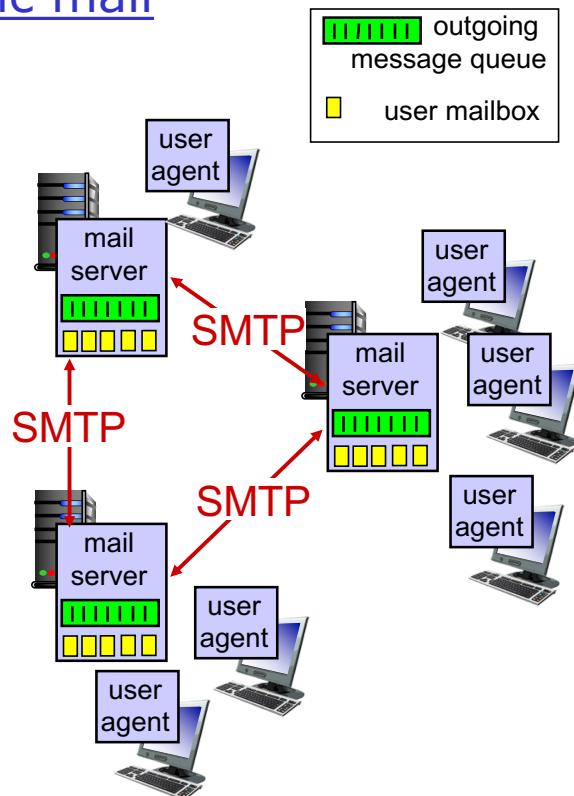
Electronic mail

Three major components:

- ❑ user agents
- ❑ mail servers
- ❑ simple mail transfer protocol (SMTP)

User Agent

- ❑ a.k.a. “mail reader”
- ❑ composing, editing, reading mail messages
- ❑ e.g., Microsoft Outlook, Apple Mail
- ❑ outgoing, incoming messages stored on server

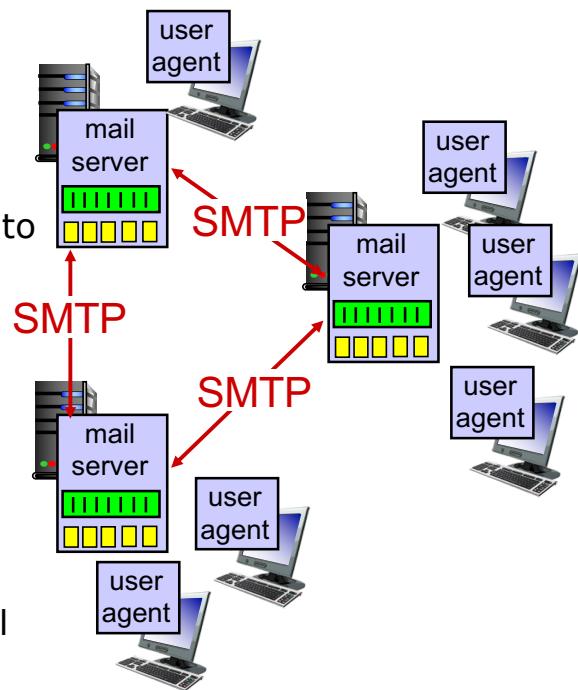


ELEC 331 45

Electronic mail: mail servers

Mail Servers

- ❑ **mailbox** contains incoming messages for user
- ❑ **message queue** of outgoing (to be sent) mail messages
- ❑ **SMTP protocol** between mail servers to send email messages
 - client side: sending mail server
 - server side: receiving mail server



ELEC 331 46

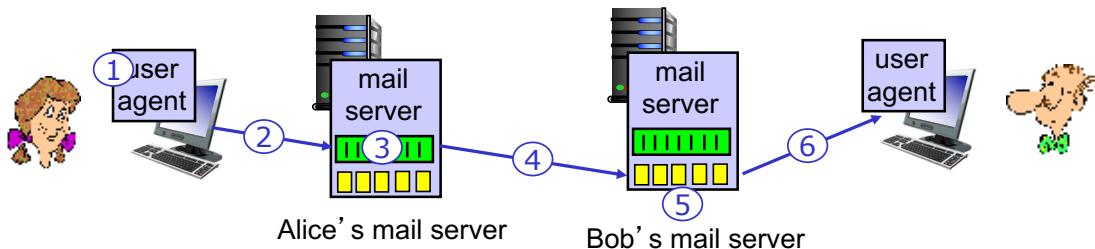
Electronic mail: SMTP [RFC 5321]

- ❑ uses TCP to reliably transfer email message from client to server, port 25
- ❑ direct transfer: sending server to receiving server
- ❑ three phases of transfer
 - handshaking (greeting)
 - transfer of messages
 - closure
- ❑ command/response interaction
 - commands: ASCII text
 - response: status code and phrase
- ❑ messages must be in 7-bit ASCII

ELEC 331 47

Scenario: Alice sends message to Bob

- 1) Alice invokes her user agent (UA) to compose message and “to” bob@someschool.edu
- 2) Alice’s UA sends message to her mail server; message placed in message queue
- 3) Client side of SMTP opens TCP connection with Bob’s mail server
- 4) SMTP client sends Alice’s message over the TCP connection
- 5) Bob’s mail server places the message in Bob’s mailbox
- 6) Bob invokes his user agent to read message



ELEC 331 48

Sample SMTP interaction

```
S: 220 hamburger.edu
C: HELO crepes.fr
S: 250 Hello crepes.fr, pleased to meet you
C: MAIL FROM: <alice@crepes.fr>
S: 250 alice@crepes.fr... Sender ok
C: RCPT TO: <bob@hamburger.edu>
S: 250 bob@hamburger.edu ... Recipient ok
C: DATA
S: 354 Enter mail, end with "." on a line by itself
C: Do you like ketchup?
C: How about pickles?
C: .
S: 250 Message accepted for delivery
C: QUIT
S: 221 hamburger.edu closing connection
```

ELEC 331 49

Try SMTP interaction for yourself:

- ❑ **telnet servername 25**
where **servername** is the name of the *local* mail server.
- ❑ see 220 reply from server
- ❑ enter HELO, MAIL FROM, RCPT TO, DATA, QUIT commands

above lets you send email without using email client (reader)

ELEC 331 50

SMTP: final words

- ❑ SMTP uses persistent connections
 - ❑ SMTP requires message (header & body) to be in 7-bit ASCII. HTTP does not have this restriction.
 - ❑ SMTP server uses CRLF.CRLF to determine end of message
- Comparison with HTTP:
- ❑ Both use persistent connections.
 - ❑ HTTP: **pull** protocol
 - ❑ SMTP: **push** protocol
 - ❑ both have ASCII command/response interaction, status codes
 - ❑ HTTP: each object encapsulated in its own response message
 - ❑ SMTP: multiple objects sent in single multipart message

ELEC 331 51

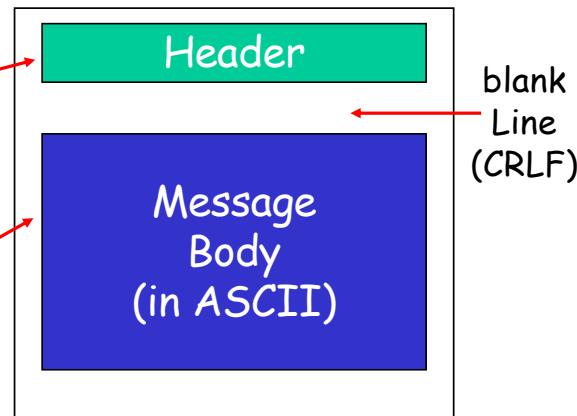
Mail message format

SMTP: protocol for exchanging email messages

RFC 5322: standard for text message format:

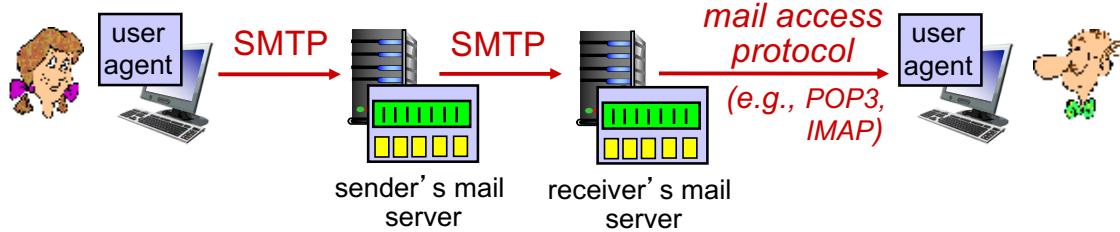
- ❑ **header lines**, e.g.,
 - From:
 - To:
 - Subject:
- different from SMTP MAIL FROM, RCPT TO commands!*

- ❑ **body**
 - the “message”, ASCII characters only



ELEC 331 52

Mail access protocols



- **SMTP**: delivery/storage to receiver's server
- Mail access protocol: retrieval from server
 - **POP3**: Post Office Protocol [RFC 1939]
 - authorization (agent <-->server) and download
 - **IMAP**: Internet Mail Access Protocol [RFC 3501]
 - more features, including manipulation of stored msgs on server
 - **HTTP**: e.g., gmail, hotmail
 - UA communicates (i.e., sends/retrieves e-mail msgs) with its mail server using HTTP GET and POST.

ELEC 331 53

POP3 protocol

UA opens TCP connection to server on port 110.

authorization phase →

- client commands:

- **user**: declare username
 - **pass**: password

- server responses

- **+OK**
 - **-ERR**

transaction phase, client: →

- **list**: list message numbers
- **retr**: retrieve message by number
- **delete**: delete

update phase, client: →

- **quit**

```
S: +OK POP3 server ready
C: user bob
S: +OK
C: pass hungry
S: +OK user successfully logged on

C: list
S: 1 498
S: 2 912
S: .
C: retr 1
S: <message 1 contents>
S: .
C: dele 1
C: retr 2
S: <message 2 contents>
S: .
C: dele 2
C: quit
S: +OK POP3 server signing off
```

ELEC 331 54

POP3 (more) and IMAP

More about POP3

- ❑ Previous example uses “download-and-delete” mode
- ❑ Bob cannot re-read e-mail if he changes client
- ❑ “Download-and-keep”: copies of messages on different clients
- ❑ POP3 is stateless across sessions

IMAP

- ❑ Keep all messages in one place: the server
- ❑ Allow user to organize and maintain msgs in **remote folders**
- ❑ user can **search** remote folders for msg matching some criteria
- ❑ keep user state across sessions
 - names of folders and mappings between msg IDs and folder name
- ❑ Permit user agent obtains components of message

ELEC 331 55

Chapter 2: Application layer

- ❑ 2.1 Principles of network applications
- ❑ 2.2 Web and HTTP
- ❑ 2.3 Electronic Mail
 - SMTP, POP3, IMAP
- ❑ 2.4 DNS
- ❑ 2.5 P2P applications
- ❑ 2.6 video streaming and content distribution networks
- ❑ 2.7 Socket programming with UDP and TCP

ELEC 331 56

DNS: Domain Name System (RFC 1034, 1035)

People: many identifiers:

- name, social insurance number, passport #

Internet hosts, routers:

- **hostname**, e.g., www.yahoo.com - used by humans
- **IP address** (32 bit) - e.g., 121.7.106.83, used for addressing datagrams

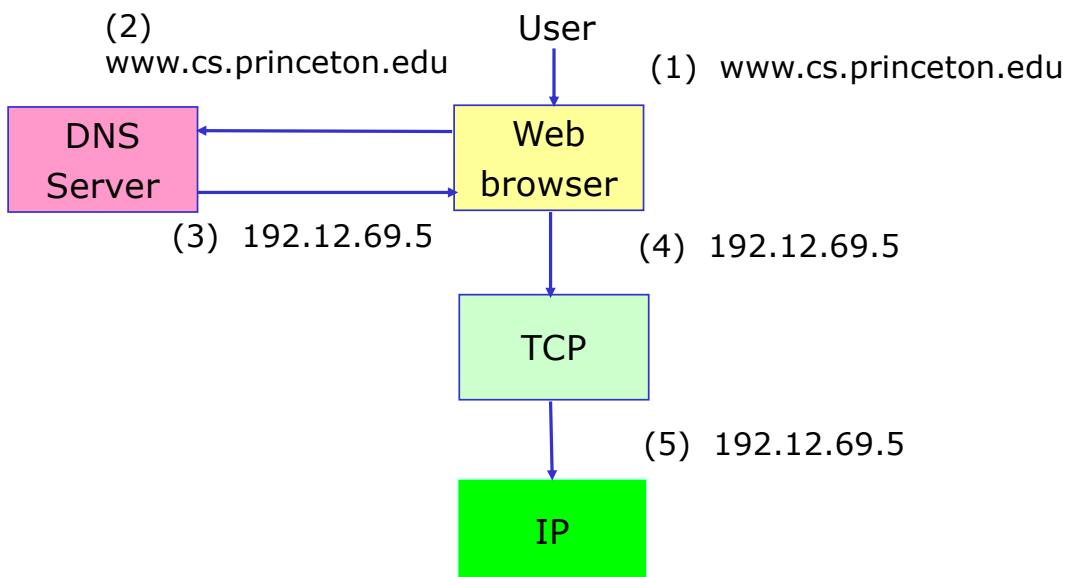
Q: Mapping between hostname and IP addresses?

Domain Name System:

- *distributed database* implemented in hierarchy of many *DNS servers*
- *application-layer protocol* : allow host, routers, DNS servers to communicate to *resolve* names (address/name translation)
 - note: core Internet function, implemented as application-layer protocol
 - complexity at network edge
- *Runs over UDP, Uses port 53*

ELEC 331 57

DNS example



- Names translated into addresses, where the numbers (1)-(5) show the sequence of steps in the process.

ELEC 331 58

DNS services

1. Hostname to IP address translation
 - ❑ e.g. hostname = mach4.ece.ubc.ca, IP Addr = 137.82.57.104
2. Host aliasing
 - ❑ alias name and canonical hostname (real name) translation
 - ❑ alias name = ca.yahoo.com, canonical hostname = media-router-fp1.media.yahoo.com
3. Mail server aliasing
 - ❑ alias name = ece.ubc.ca, canonical name of mail server = esva.mail-relay.ubc.ca
4. Load distribution
 - ❑ replicated Web servers: many IP addresses correspond to one canonical name
 - ❑ hostname = www.yahoo.com, set of IP addresses: 206.190.39.43, 206.190.39.42.

ELEC 331 59

DNS

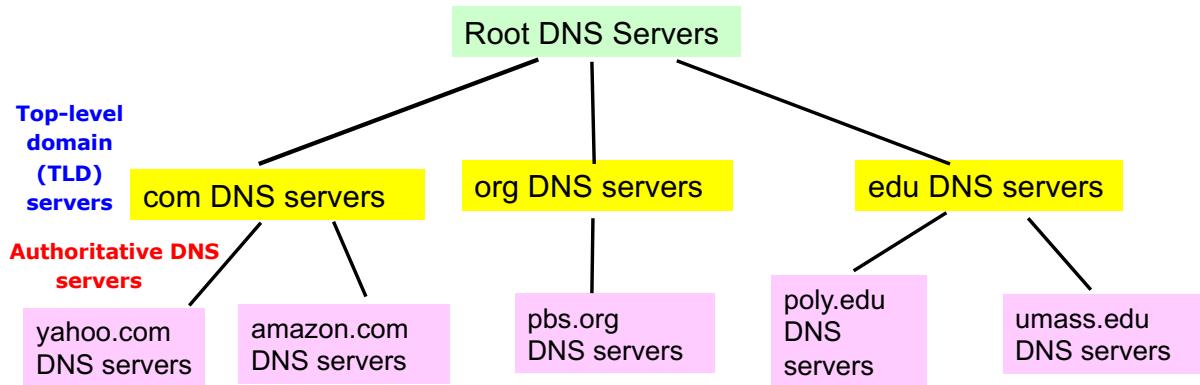
Why not using a centralized DNS?

- ❑ single point of failure
- ❑ traffic volume
- ❑ distant centralized database
- ❑ maintenance

doesn't scale!

ELEC 331 60

Distributed, hierarchical database



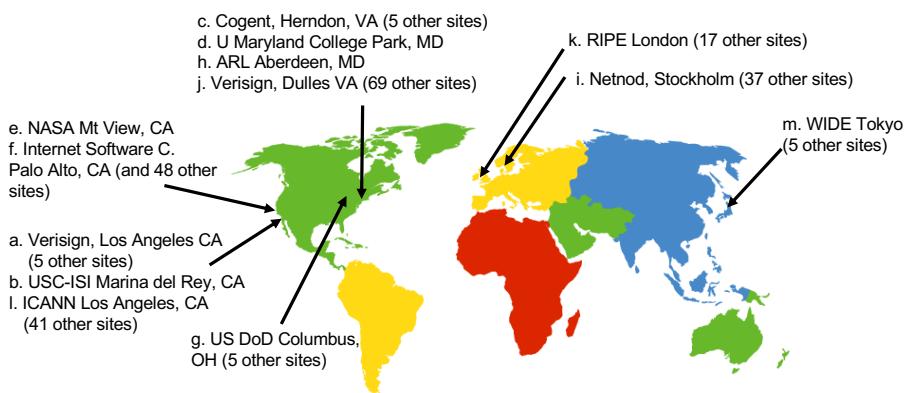
Client wants IP address for www.amazon.com; 1st approximation

- client queries a root server to find com DNS server
- client queries com DNS server to get amazon.com DNS server
- client queries amazon.com DNS server to get IP address for www.amazon.com

ELEC 331 61

Root DNS servers

- Contacted by local name server that cannot resolve name
- Provide IP addresses of the TLD servers
- Over 400 root DNS servers scattered all over the world for both security and reliability purposes
- These root DNS servers are managed by 13 different organizations
- <http://www.root-servers.org/>



ELEC 331 62

TLD and authoritative servers

- ❑ **Top-level domain (TLD) servers:** responsible for com, org, net, edu, etc, and all top-level country domains uk, fr, ca, jp.
 - Provide IP addresses for authoritative DNS servers
 - [Verisign Global Registry Services](#) maintains TLD servers for .com top-level domain
 - [Educause](#) maintains TLD servers for .edu top-level domain
- ❑ **Authoritative DNS servers:** organization's DNS servers, providing authoritative hostname to IP mappings for organization's servers (e.g., Web and mail).
 - Provide publicly accessible DNS records
 - Can be maintained by organization or service provider
 - In our department, it is the “dns.ece.ubc.ca”

ELEC 331 63

Local DNS server

- ❑ Does not strictly belong to hierarchy
- ❑ Each ISP (residential ISP, company, university) has one.
 - Also called “[default name server](#)”
- ❑ When a host makes a DNS query, query is sent to its local DNS server
 - Acts as a proxy, forwards query into hierarchy.
 - e.g., in our department, the local name servers are:
dns1.ece.ubc.ca, dns2.ece.ubc.ca

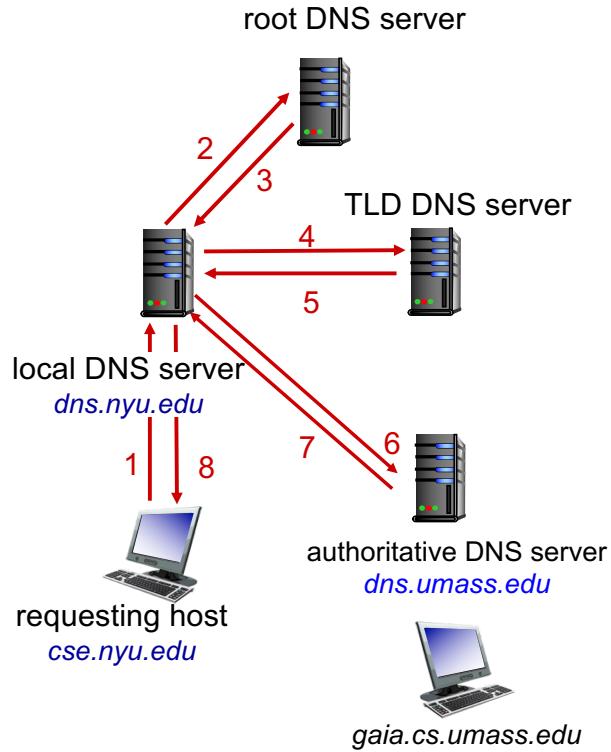
ELEC 331 64

DNS name resolution example

- ❑ Host at cse.nyu.edu wants IP address for gaia.cs.umass.edu

iterated query:

- ❑ contacted server replies with name of server to contact
- ❑ “I don’t know this name, but ask this server”

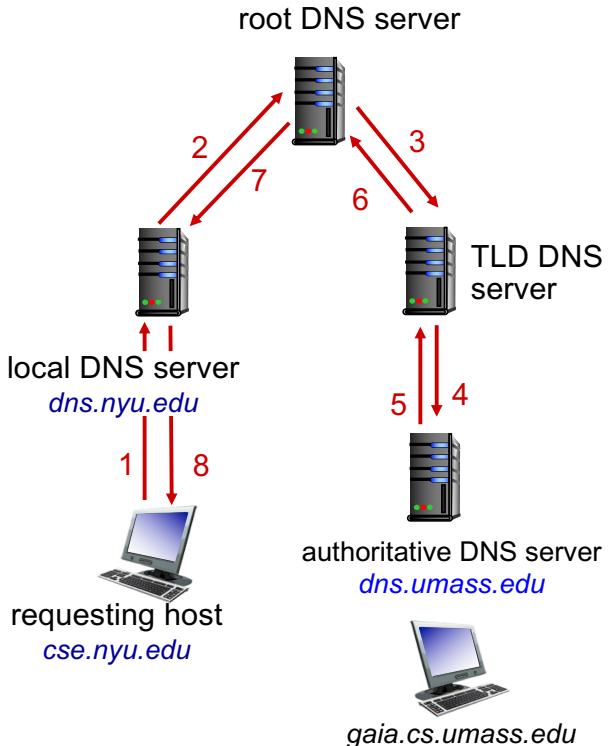


ELEC 331 65

DNS name resolution example

Recursive query:

- ❑ puts burden of name resolution on contacted name server
- ❑ heavy load at upper levels of hierarchy



- ❑ In practice, the query from requesting host to local DNS server is **recursive**, and the remaining queries are **iterative**.

ELEC 331 66

DNS: Caching and updating records

- ❑ once (any) name server learns mapping, it *caches* mapping
 - cache entries timeout (disappear) after some time (TTL) (e.g., two days).
 - TLD servers typically cached in local name servers.
 - Thus root name servers not often visited.
- ❑ cached entries may be **out-of-date** (best effort name-to-address translation!)
 - if name host changes IP address, may not be known Internet-wide until all TTLs expire
- ❑ update/notify mechanisms under design by IETF
 - RFC 2136
 - <http://www.ietf.org/html.charters/dnsind-charter.html>

ELEC 331 67

DNS records

DNS: distributed database storing **resource records (RRs)**

RR format: (**name**, **value**, **type**, **TTL**)

- ❑ Type=**A**
 - **name** is hostname
 - **value** is IP address
 - (relay1.bar.foo.com, 145.37.93.2, A)
- ❑ Type=**NS**
 - **name** is domain (e.g., foo.com)
 - **value** is hostname of authoritative DNS server for this domain
 - (foo.com, dns.foo.com, NS)
- ❑ Type=**CNAME**
 - **name** is alias name for some “canonical” (the real) name
 - **value** is canonical name
 - (foo.com, relay1.bar.foo.com, CNAME)
- ❑ Type=**MX**
 - **value** is name of mailserver associated with **name**
 - (foo.com, mail.bar.foo.com, MX)

ELEC 331 68

DNS records example

- ❑ Sample records maintained by root server

```
<edu, a3.nstld.com, NS>
<a3.nstld.com, 192.5.6.32, A>
<com, a.gtld-servers.net, NS>
<a.gtld-servers.net, 192.5.6.30, A>
...
...
```

- ❑ Records maintained by a3.nstld.com server for .edu domain

```
<princeton.edu, dns.princeton.edu, NS>
<dns.princeton.edu, 128.112.125.15, A>
...
...
```

- ❑ Records maintained by authoritative server dns.princeton.edu

```
<cs.princeton.edu, dns1.cs.princeton.edu, NS>
<dns1.cs.princeton.edu, 128.112.136.10, A>
...
...
```

ELEC 331 69

DNS records example (cont.)

- ❑ Sample records maintained by authoritative server dns1.cs.princeton.edu

```
<penguins.cs.princeton.edu, 128.112.155.166, A>
<www.cs.princeton.edu, coreweb.cs.princeton.edu, CNAME>
<coreweb.cs.princeton.edu, 128.112.136.35, A>
<cs.princeton.edu, mail.cs.princeton.edu, MX>
<mail.cs.princeton.edu, 128.112.136.72, A>
...
...
```

- ❑ Local name server maintains the cache and also records for root servers

```
<'root', a.root-servers.net, NS>
<a.root-servers.net, 198.41.0.4, A>
...
...
```

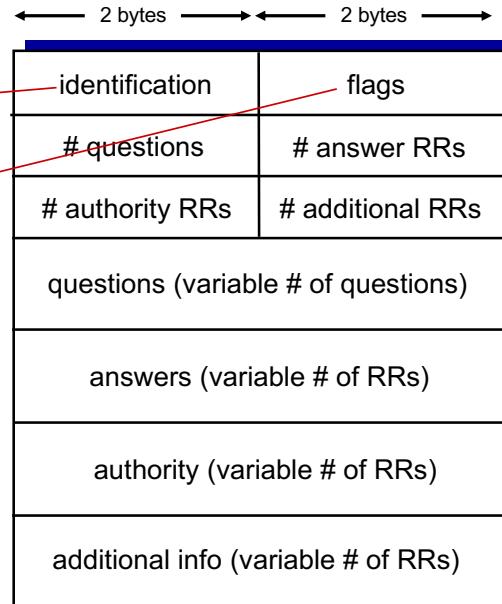
ELEC 331 70

DNS protocol, messages

DNS protocol: *query* and *reply* messages, both with same
message format

Message Header

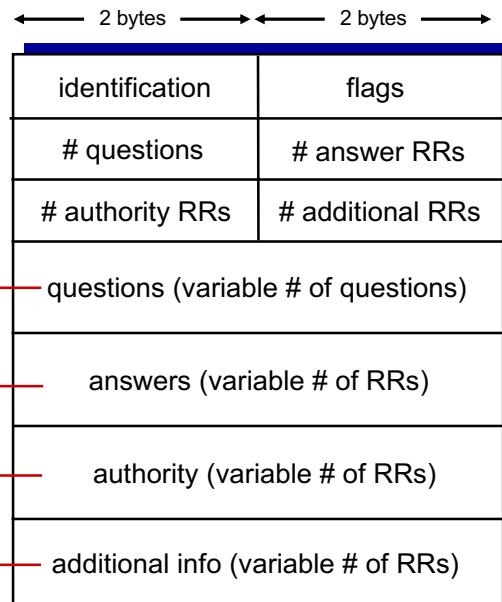
- ❑ **identification:** 16 bit # for query, reply to query uses same #
- ❑ **flags:**
 - query (0) or reply (1)
 - reply is authoritative (1)
 - recursion desired: set by client
 - recursion available: set by server



ELEC 331 71

DNS protocol, messages

- name, type fields for a query
- RRs in response to query
- records for authoritative servers
- additional “helpful” info that may be used



ELEC 331 72

Inserting records into DNS

- ❑ Example: new startup “Network Utopia”
- ❑ Register name networkutopia.com at a DNS **registrar** (e.g., Network Solutions)
 - Provide names and IP addresses of authoritative name server (primary and secondary)
 - Registrar inserts two RRs into the com TLD server:
(networkutopia.com, dns1.networkutopia.com, NS)
(dns1.networkutopia.com, 212.212.212.1, A)
- ❑ create authoritative server Type A record for www.networkutopia.com and Type MX record for mail.networkutopia.com
 - (www.networkutopia.com, 192.10.232.2, A)
 - (networkutopia.com, mail.networkutopia.com, MX)
 - (mail.networkutopia.com, 192.10.232.3, A)

ELEC 331 73

Attacking DNS

DDoS attacks

- ❑ Bombard root servers with traffic
 - Not successful to date
 - packet filters
 - Local DNS servers cache IP address of TLD servers, allowing to bypass root server
- ❑ Bombard TLD servers
 - Potentially more dangerous

Redirect attacks

- ❑ Man-in-middle
 - Intercept queries and returns bogus replies
- ❑ DNS poisoning
 - Send bogus replies to DNS server, which caches
- ❑ Difficult to implement

Exploit DNS for reflection attack

- ❑ Send queries with spoofed source address: target IP
- ❑ Size of reply message > > Size of query message

ELEC 331 74

Chapter 2: Application layer

- ❑ 2.1 Principles of network applications
 - ❑ app architectures
 - ❑ app requirements
- ❑ 2.2 Web and HTTP
- ❑ 2.3 Electronic Mail
 - ❑ SMTP, POP3, IMAP
- ❑ 2.4 DNS
- ❑ 2.5 P2P applications
- ❑ 2.6 Video streaming and content distribution networks
- ❑ 2.7 Socket programming with UDP and TCP

ELEC 331 75

Peer-to-Peer applications

1. File Distribution
 - ❑ Distribute a file from a single source to a large number of peers
 - ❑ BitTorrent protocol
2. Searching for Information
 - ❑ Support search and update operation
 - ❑ Distributed Hash Tables (DHTs)
3. Skype: P2P Internet telephony application

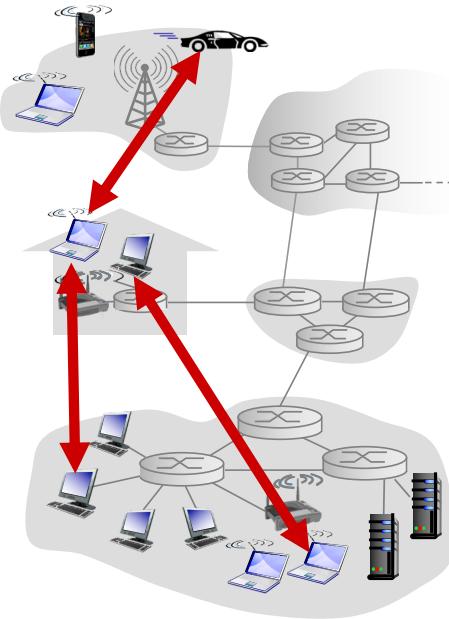
ELEC 331 76

P2P architecture

- ❑ no always-on server
- ❑ arbitrary end systems directly communicate
- ❑ peers are intermittently connected and change IP addresses

Examples:

- ❑ File distribution (BitTorrent)
- ❑ Streaming (KanKan)
- ❑ VoIP (Skype)

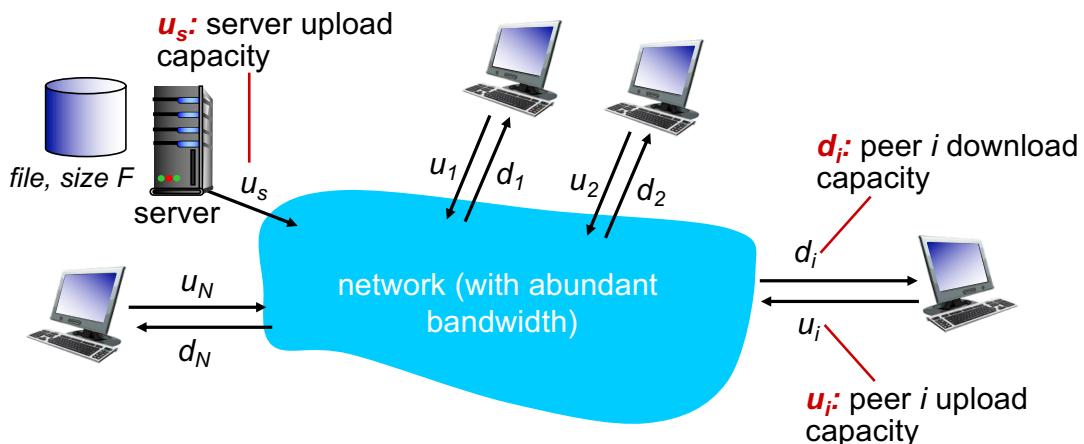


ELEC 331 77

File distribution: client-server vs P2P

Question: how much time to distribute file (size F bits) from one server to N peers?

- ❑ peer upload/download capacity is limited resource.



ELEC 331 78

File distribution time: client-server

- server transmission: must sequentially send (upload) N file copies:

 - time to send one copy: F/u_s
 - time to send N copies: NF/u_s

- client: each client must download file copy

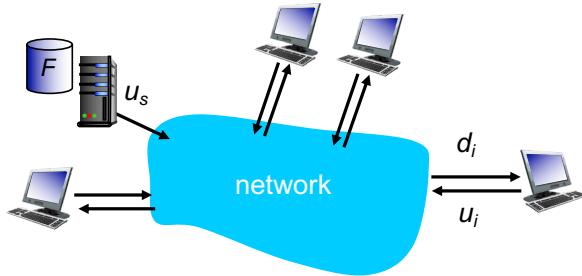
 - d_{min} = min client download rate
 - download time for client with bottleneck link d_{min} : F/d_{min}

time to distribute F bits to N clients using client-server approach

$$D_{cs} \geq \max\{ NF/u_s, F/d_{min} \}$$

increases linearly in N

ELEC 331 79



File distribution time: P2P

- server transmission: must upload at least one copy:

 - time to send one copy: F/u_s

- client: each client must download file copy

 - min client download time: F/d_{min}

- clients: as aggregate must download NF bits

 - max upload rate (limiting max download rate) is $u_s + \sum u_i$

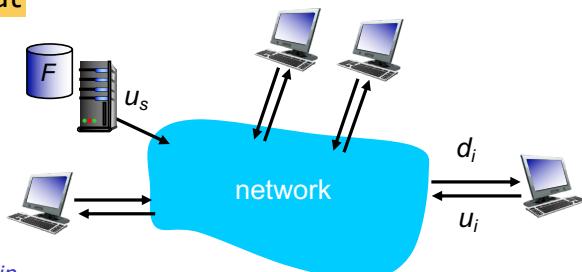
time to distribute F bits to N clients using P2P approach

$$D_{P2P} \geq \max\{F/u_s, F/d_{min}, NF/(u_s + \sum u_i)\}$$

increases linearly in N ...

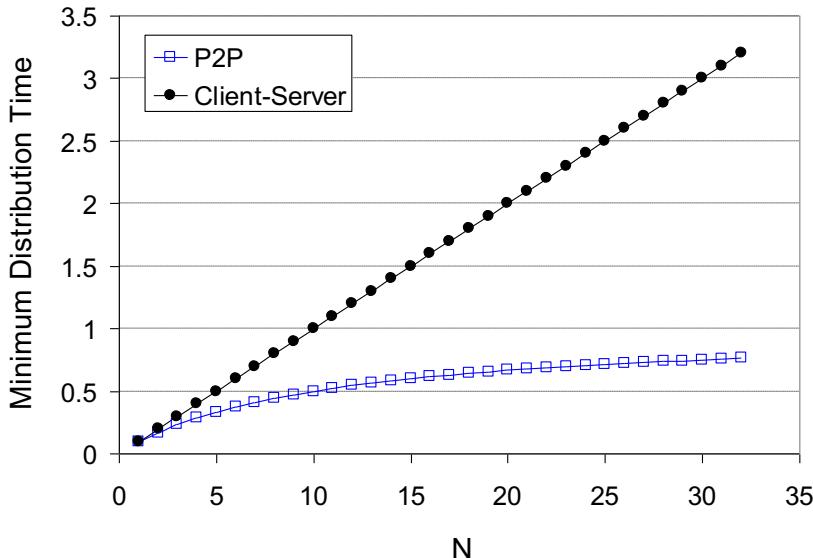
... but so does this, as each peer brings service capacity

ELEC 331 80



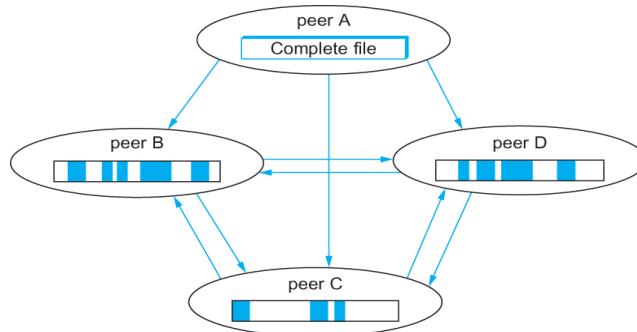
Client-server vs. P2P: example

Client upload rate = u , $F/u = 1$ hour, $u_s = 10u$, $d_{min} \geq u_s$



- ❑ Peers being consumers as well as redistributors of bits.

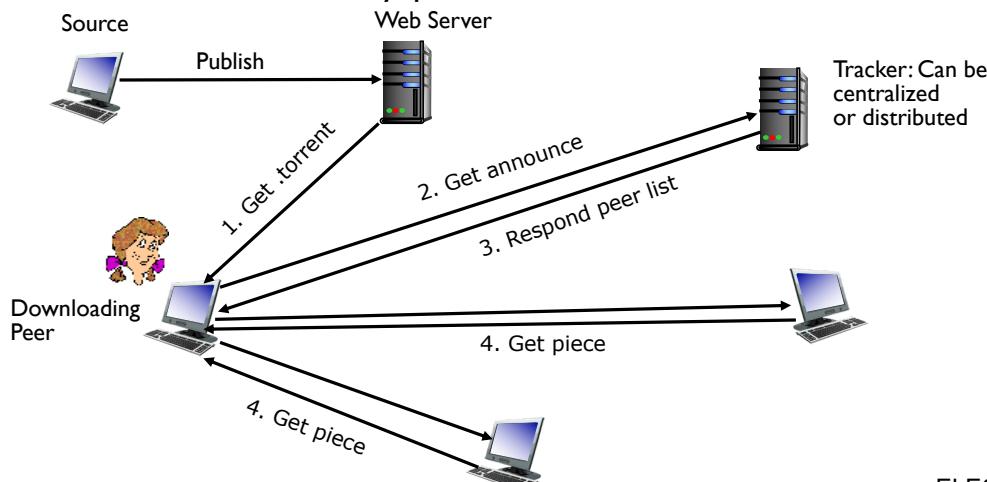
BitTorrent basics



- ❑ File is divided into fixed size (e.g., 256 Kbytes) chunks
- ❑ Peers in torrent send/receive file chunks
- ❑ Peers in a BitTorrent swarm download from other peers that may not yet have the complete file
- ❑ **Replication** is a natural side-effect
- ❑ As soon as a peer downloads a particular piece, it becomes another source for that piece

P2P file distribution: BitTorrent

- ❑ To share a file, a peer first creates a “.torrent” file
- ❑ .torrent file: contains metadata about the file to be shared (e.g., file name, file size, size of each piece/chunk, SHA-1 hash code of each piece, URL of tracker)
- ❑ Torrent file is usually published to some well-known website.

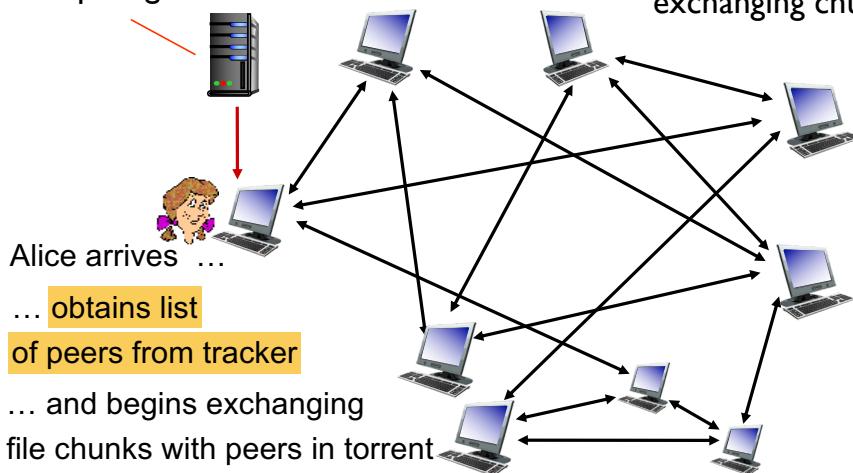


ELEC 331 83

P2P file distribution: BitTorrent

tracker: tracks peers
participating in torrent

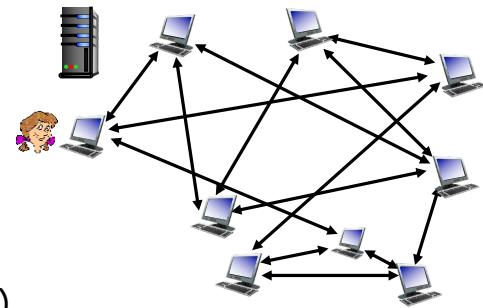
torrent: group of peers
exchanging chunks of a file



ELEC 331 84

BitTorrent Schematic

- ❑ peer joining torrent:
 - has no chunks, but will accumulate them over time from other peers
 - registers with tracker to get list of peers, connects to subset of peers ("neighbors")
- ❑ while downloading, peer uploads chunks to other peers
- ❑ peer may change peers with whom it exchanges chunks
- ❑ **churn**: refers to the phenomenon that peers may dynamically join and leave the system at well
- ❑ once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent



ELEC 331 85

Piece Selection

- ❑ At any given time, different peers have different subsets of file chunks
- ❑ Periodically, Alice asks her neighbors for list of chunks that they have
- ❑ Use **random first selection** at the initial phase
 - Randomly select the first few pieces (usually around four)
- ❑ After that, use **rarest first policy**
 - Select those chunks which have the fewest repeated copies
- ❑ At the end, **end-game mode**
 - Send requests for all missing pieces to all the peers

ELEC 331 86

Peer Selection

sending chunks: tit-for-tat

- ❑ Alice sends chunks to those four peers currently supplying her data chunks at the highest rate
 - these four peers are said to be **unchoked**
 - re-evaluate top 4 every 10 secs
- ❑ every 30 secs: randomly select another peer, starts sending chunks
 - This peer is said to be **optimistically unchoked**
 - **newly chosen peer may join top 4**
- ❑ other peers are **choked** by Alice (do not receive chunks from her)

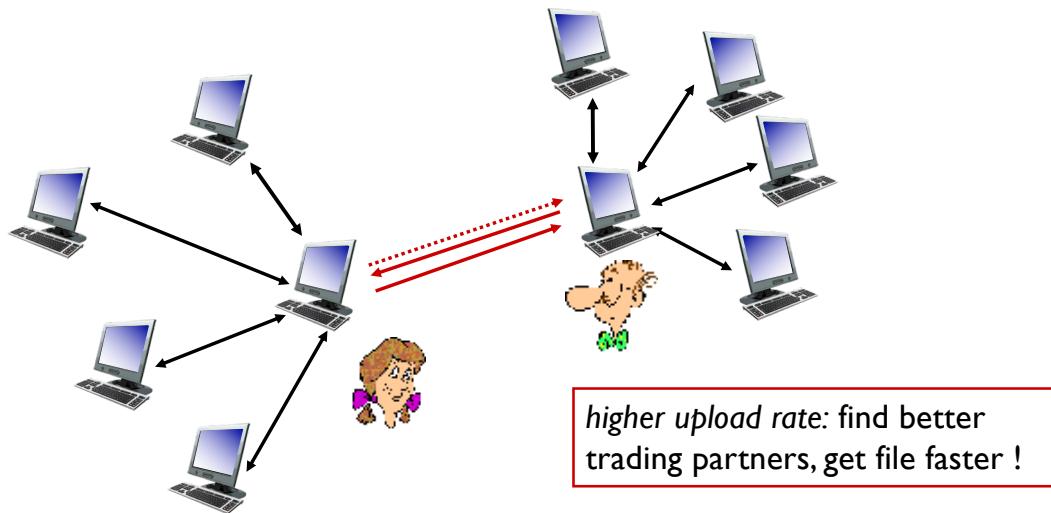
Peer Selection

Anti-snubbing algorithm

- ❑ Alice runs the anti-snubbing algorithm whenever she is choked by all of her peers
 - A peer is snubbed if it does not receive any data in 60 sec.
- ❑ Alice can do better to run optimistically unchoking more often to explore more peers who are willing to cooperate.
- ❑ Thus, the anti-snubbing algorithm stops uploading to peers selected tit-for-tat so that optimistically unchoking can be performed more often.

BitTorrent: Tit-for-tat

- (1) Alice “optimistically unchoke” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



ELEC 331 89

Chapter 2: Application layer

- ❑ 2.1 Principles of network applications
 - ❑ app architectures
 - ❑ app requirements
- ❑ 2.2 Web and HTTP
- ❑ 2.3 Electronic Mail
 - ❑ SMTP, POP3, IMAP
- ❑ 2.4 DNS
- ❑ 2.5 P2P applications
- ❑ 2.6 Video streaming and content distribution networks
- ❑ 2.7 Socket programming with UDP and TCP

ELEC 331 90

Video Streaming and CDNs: context

- ❑ video traffic: major consumer of Internet bandwidth
 - Netflix, YouTube: 37%, 16% of downstream residential ISP traffic
 - ~1B YouTube users, ~75M Netflix users
- ❑ challenge: scale - how to reach ~1B users?
 - single mega-video server won't work (why?)
- ❑ challenge: heterogeneity
 - different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)
- ❑ solution: distributed, application-level infrastructure

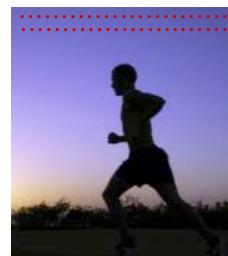


ELEC 331 91

Multimedia: video

- ❑ video: sequence of images displayed at constant rate
 - e.g., 24 images/sec
- ❑ digital image: array of pixels
 - each pixel represented by bits
- ❑ coding: use redundancy within and between images to decrease # bits used to encode image
 - spatial (within image)
 - temporal (from one image to next)

spatial coding example: instead of sending N values of same color (all purple), send only two values: color value (purple) and number of repeated values (N)



frame i

temporal coding example: instead of sending complete frame at $i + 1$, send only differences from frame i



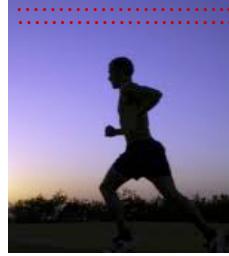
frame $i + 1$

ELEC 331 92

Multimedia: video

- ❑ CBR: (constant bit rate):
video encoding rate fixed
- ❑ VBR: (variable bit rate):
video encoding rate changes
as amount of spatial,
temporal coding changes
- ❑ examples:
 - MPEG 1 (CD-ROM) 1.5 Mbps
 - MPEG2 (DVD) 3-6 Mbps
 - MPEG4 (often used in Internet, < 1 Mbps)

spatial coding example: instead of sending N values of same color (all purple), send only two values: color value (purple) and number of repeated values (N)



frame i

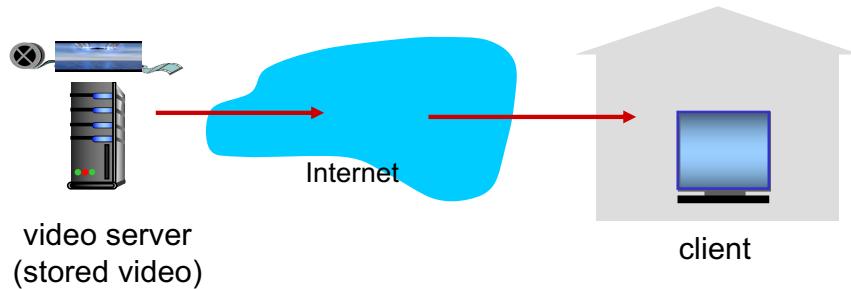


frame $i + 1$

temporal coding example: instead of sending complete frame at $i + 1$, send only differences from frame i

Streaming stored video

simple scenario:



Streaming multimedia: DASH

- **DASH:** Dynamic Adaptive Streaming over HTTP
- A video is encoded into several different versions, each having a different bit rate (e.g., 300 kbps, 1 Mbps, 3 Mbps)
- Each video version is stored in server, each with a different URL
- Server has a *manifest file* which provides a URL for each version along with its bit rate
- **client:**
 - Select one chunk at a time by specifying a URL and a byte range in HTTP GET request message for each chunk.
 - periodically measures server-to-client bandwidth by **rate determination algorithm**
 - can choose different coding rates at different points in time (depending on available bandwidth at time)

ELEC 331 95

Streaming multimedia: DASH

- **DASH:** Dynamic Adaptive Streaming over HTTP
- “intelligence” at client: client determines
 - **when** to request chunk (so that buffer starvation, or overflow does not occur)
 - **what encoding rate** to request (higher quality when more bandwidth available)
 - **where** to request chunk (can request from URL server that is “close” to client or has high available bandwidth)

ELEC 331 96

Content distribution networks

- **challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?
- **option 1:** single, large “mega-server”
 - single point of failure
 - point of network congestion
 - long path to distant clients
 - multiple copies of video sent over outgoing link
-quite simply: this solution **doesn't scale**

ELEC 331 97

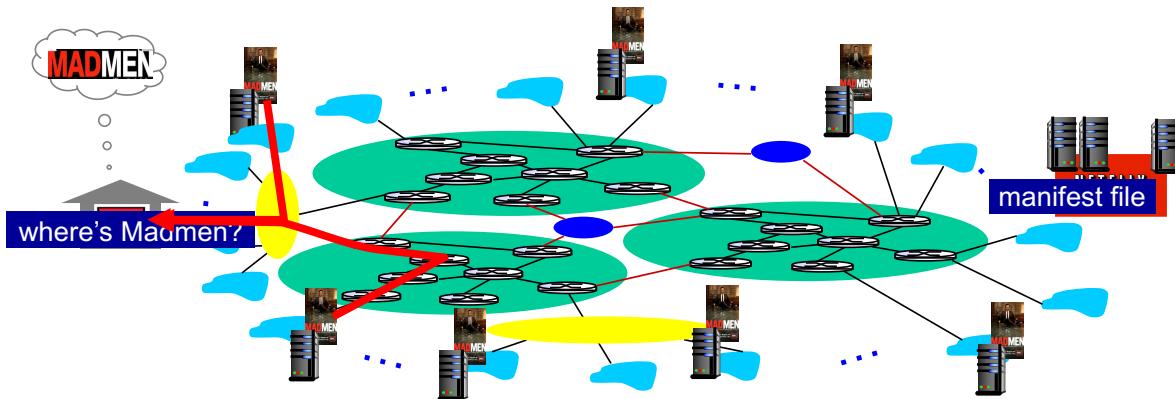
Content distribution networks

- **challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?
- **option 2:** store/serve multiple copies of videos at multiple geographically distributed sites (**CDN**)
 - *enter deep*: push CDN servers deep into many access networks
 - close to users
 - used by Akamai, 1700 locations
 - *bring home*: smaller number (10's) of larger clusters in IXPs near (but not within) access networks
 - used by Limelight
 - lower maintenance and management overhead

ELEC 331 98

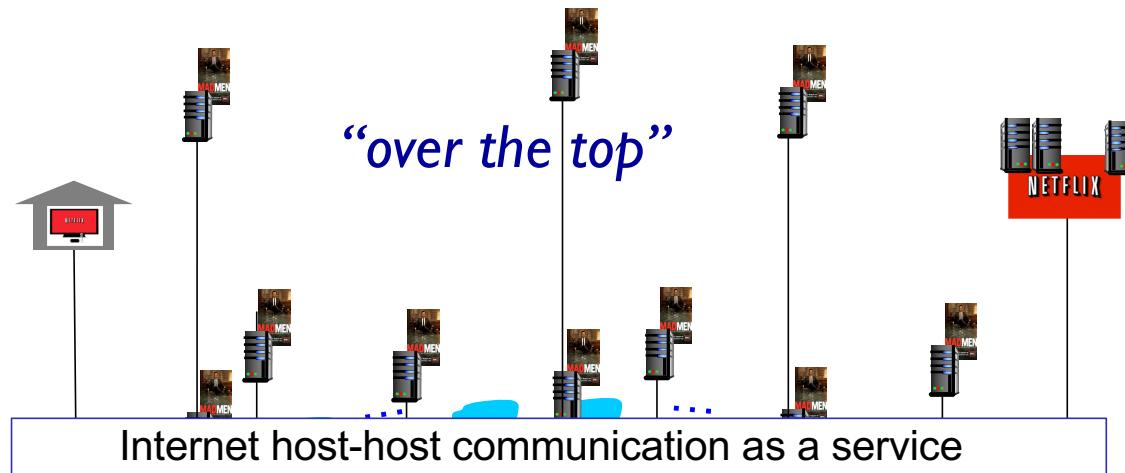
Content Distribution Networks (CDNs)

- CDN: stores copies of content at CDN nodes
 - e.g. Netflix stores copies of MadMen
- subscriber requests content from CDN
 - redirected to nearby CDN server cluster, retrieves content
 - may choose different server if network path congested



ELEC 331 99

Content Distribution Networks (CDNs)



OTT challenges: coping with a congested Internet

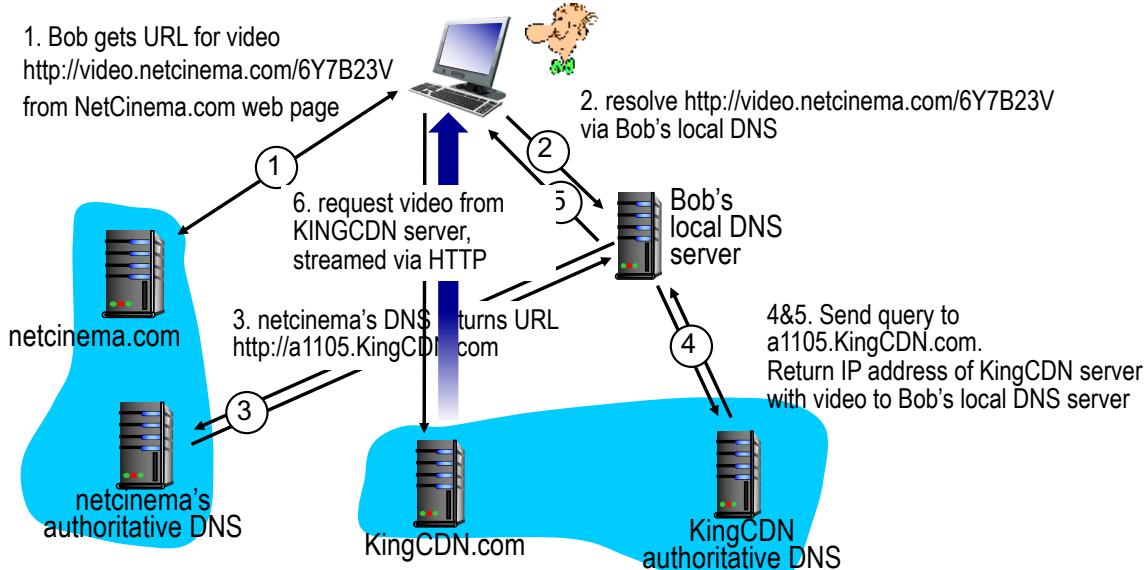
- from which CDN node to retrieve content?
- viewer behavior in presence of congestion?
- what content to place in which CDN node?

more .. in Chapter 9

ELEC 331 100

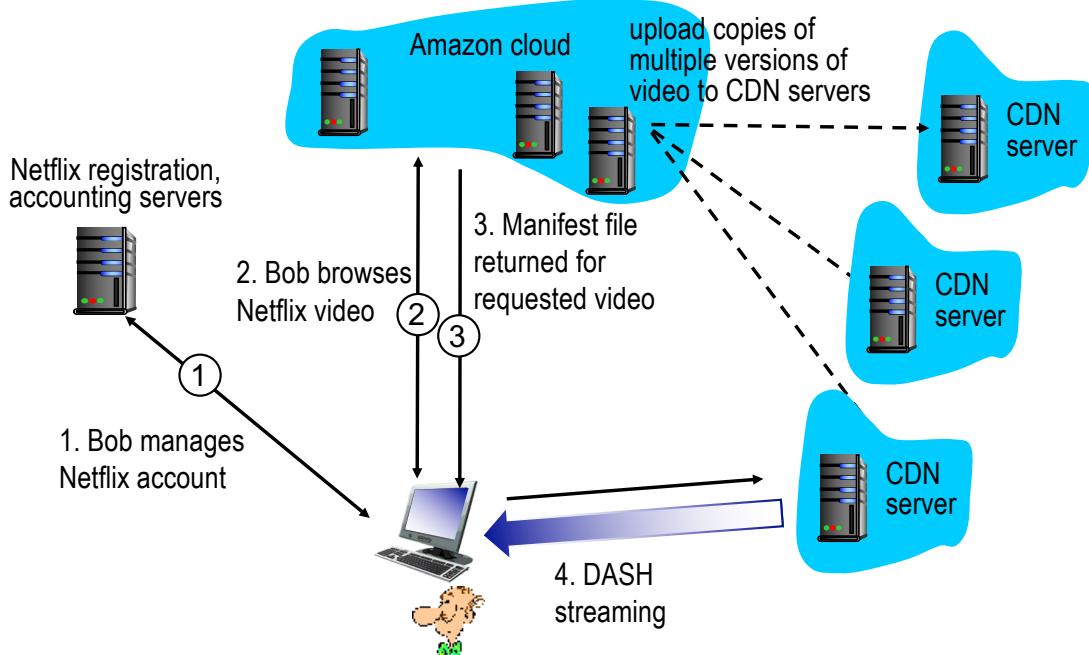
CDN content access: a closer look

Bob (client) requests video <http://video.netcinema.com/6Y7B23V>



ELEC 331 101

Case study: Netflix



ELEC 331 102

Chapter 2: Application layer

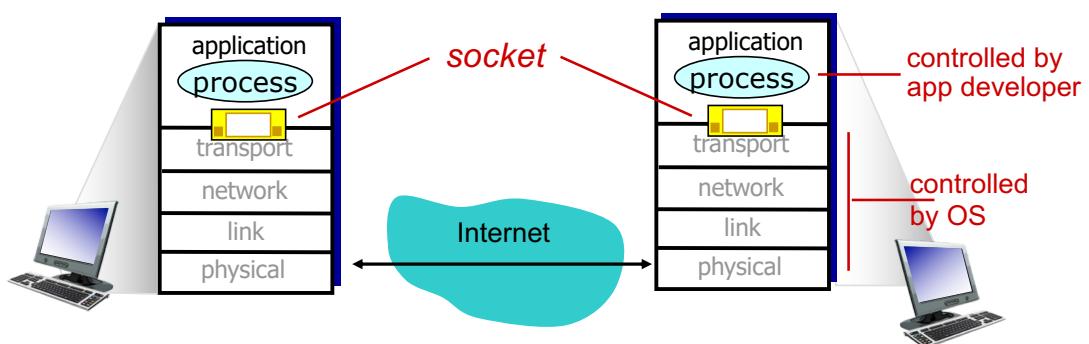
- ❑ 2.1 Principles of network applications
- ❑ 2.2 Web and HTTP
- ❑ 2.3 Electronic Mail
 - ❑ SMTP, POP3, IMAP
- ❑ 2.4 DNS
- ❑ 2.5 P2P applications
- ❑ 2.6 Video streaming and content distribution networks
- ❑ 2.7 **Socket programming with UDP and TCP**

ELEC 331 103

Socket programming

Goal: learn how to build client/server applications that communicate using sockets

Socket: door between application process and end-end-transport protocol



ELEC 331 104

Socket programming

Two socket types for two transport services:

- ❑ **UDP**: unreliable datagram
- ❑ **TCP**: reliable, byte stream-oriented

Application Example:

1. Client reads a line of characters (data) from its keyboard and sends the data to the server.
2. The server receives the data and converts characters to uppercase.
3. **The server sends the modified data to the client.**
4. The client receives the modified data and displays the line on its screen.

ELEC 331 105

Socket programming *with UDP*

UDP: no “connection” between client and server

- ❑ no handshaking
- ❑ sender explicitly attaches IP destination address and destination port number to each packet
- ❑ Receiver extracts sender (i.e., source)’s IP address, source port number from received packet

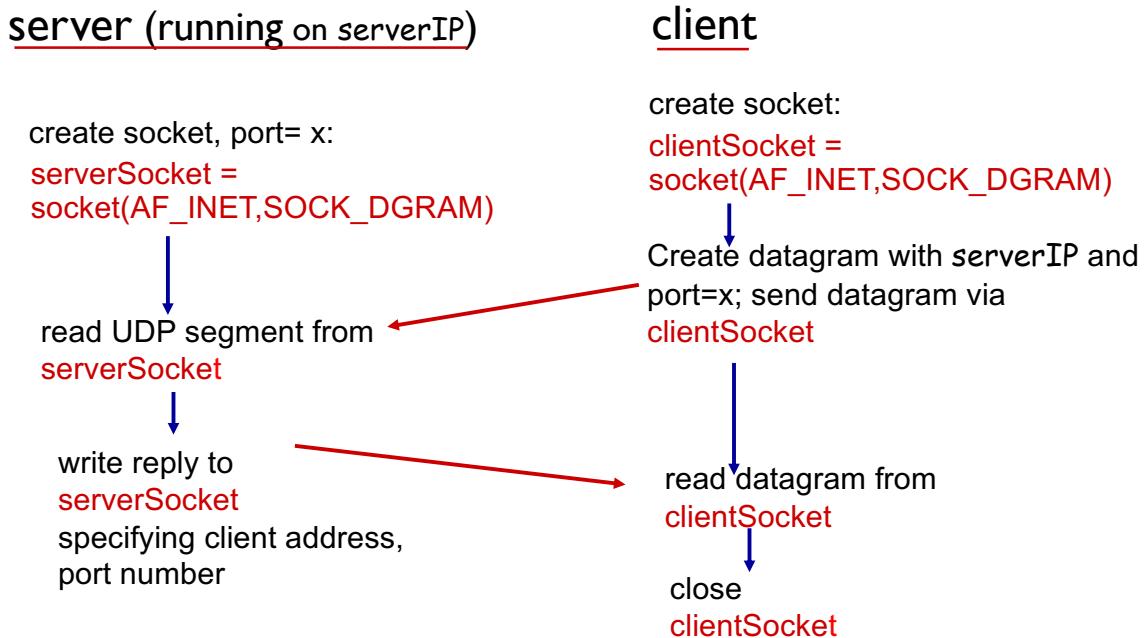
UDP: transmitted data **may be lost or received out-of-order**

Application viewpoint:

- ❑ UDP provides unreliable transfer of groups of bytes (“datagrams”) between client and server

ELEC 331 106

Client-server socket interaction: UDP



ELEC 331 107

Example app: UDP client

Python UDPClient

```
include Python's socket  
library → from socket import *  
  
serverName = 'hostname'  
serverPort = 12000  
  
create UDP socket for → clientSocket = socket(AF_INET,  
server  
get user keyboard  
input → message = raw_input('Input lowercase sentence: ')  
Attach server name, port to  
message; send into socket → clientSocket.sendto(message.encode(),  
(serverName, serverPort))  
  
read reply characters from → modifiedMessage, serverAddress =  
socket into string  
clientSocket.recvfrom(2048)  
  
print out received string → print(modifiedMessage.decode() )  
and close socket  
clientSocket.close()
```

ELEC 331 108

Example app: UDP server

Python UDPServer

```
from socket import *
serverPort = 12000
create UDP socket → serverSocket = socket(AF_INET, SOCK_DGRAM)
assign local port number → serverSocket.bind(("", serverPort))
12000 to socket
print('The server is ready to receive')
loop forever → while True:
Read from UDP socket into → message, clientAddress = serverSocket.recvfrom(2048)
message, getting client's
address (client IP and port) → modifiedMessage = message.decode().upper()
send upper case string → serverSocket.sendto(modifiedMessage.encode(),
back to this client
                                         clientAddress)
```

ELEC 331 109

Socket programming with TCP

client must contact server

- ❑ server process must first be running
- ❑ server must have created socket (door) that welcomes client's contact

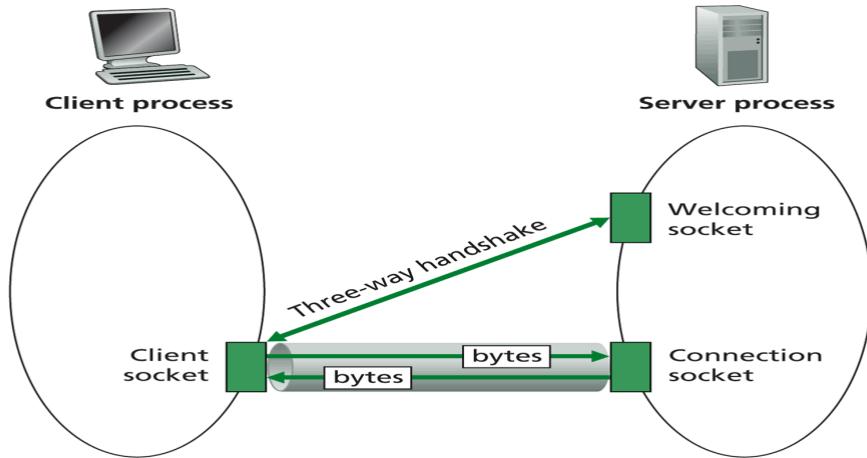
client contacts server by:

- ❑ Creating TCP socket, specifying IP address, port number of server process
- ❑ **when client creates socket:** client TCP establishes connection to server TCP

- ❑ when contacted by client, **server TCP creates new socket** for server process to communicate with that particular client
 - allows server to talk with multiple clients
 - source port numbers used to distinguish clients (more in Chap 3)

ELEC 331 110

Socket programming *with TCP*



- ◆ Client socket, welcoming socket, and connection socket

application viewpoint

TCP provides reliable, in-order transfer of bytes ("pipe") between client and server

ELEC 331 111

Client/server socket interaction: TCP

server (running on `serverIP`)

client

```
create socket,  
port=x, for incoming  
request:  
serverSocket = socket()
```

```
wait for incoming  
connection request  
connectionSocket =  
serverSocket.accept()
```

TCP
connection setup

```
create socket,  
connect to serverIP, port=x  
clientSocket = socket()
```

```
read request from  
connectionSocket
```

send request using
`clientSocket`

```
write reply to  
connectionSocket
```

read reply from
`clientSocket`

```
close  
connectionSocket
```

close
`clientSocket`

ELEC 331 112

Example app: TCP client

Python TCPClient

```
from socket import *
serverName = 'servername'
serverPort = 12000
create TCP socket for  
server, remote port 12000 → clientSocket = socket(AF_INET, SOCK_STREAM)
clientSocket.connect((serverName, serverPort))
sentence = raw_input('Input lowercase sentence: ')
No need to attach server name, port → clientSocket.send(sentence.encode())
modifiedSentence = clientSocket.recv(1024)
print ('From Server: ', modifiedSentence.decode())
clientSocket.close()
```

ELEC 331 113

Example app: TCP server

Python TCPServer

```
from socket import *
serverPort = 12000
create TCP welcoming  
socket → serverSocket = socket(AF_INET,SOCK_STREAM)
serverSocket.bind(('',serverPort))
server begins listening for  
incoming TCP requests → serverSocket.listen(1)
print('The server is ready to receive')
loop forever → while True:
server waits on accept()  
for incoming requests, new  
socket created on return → connectionSocket, addr = serverSocket.accept()
read bytes from socket (but  
not address as in UDP) → sentence = connectionSocket.recv(1024).decode()
capitalizedSentence = sentence.upper()
close connection to this  
client (but not welcoming  
socket) → connectionSocket.send(capitalizedSentence.encode())
connectionSocket.close()
```

ELEC 331 114

Chapter 2: summary

Our study of network apps now complete!

- ❑ Application architectures
 - client-server
 - P2P
- ❑ application service requirements
 - reliability, bandwidth, delay
- ❑ Internet transport service model
 - connection-oriented, reliable:
TCP
 - unreliable, datagrams: UDP
- ❑ specific protocols:
 - HTTP
 - SMTP
 - POP3, IMAP
 - DNS
 - P2P: BitTorrent
 - video streaming, CDNs
 - socket programming:
 - TCP, UDP sockets

ELEC 331 115

Chapter 2: summary

Most importantly: learned about *protocols*

- ❑ typical request/reply message exchange:
 - client requests info or service
 - server responds with data, status code
- ❑ message formats:
 - **headers**: fields giving info about data
 - **data**: info (payload) being communicated
- ❑ control vs. data messages
 - in-band, out-of-band
- ❑ centralized vs. decentralized
- ❑ stateless vs. stateful
- ❑ reliable vs. unreliable message transfer
- ❑ “complexity at network edge”

ELEC 331 116